



Karlsruhe University of Applied Sciences

Department of Computer Science and Business Information Systems

Business Information Systems

BACHELOR'S THESIS

Efficiency vs. Performance in Green AI: An Empirical Investigation of Energy Consumption, Carbon Footprint, and Predictive Accuracy of Classification Algorithms at Varying Dataset Sizes.

Author	Eron Qorri
Matriculation number	85383
Workplace	Hochschule Karlsruhe
First Advisor	Prof. Dr. rer. nat. Christine Preisach
Second Advisor	Prof. Dr. Reimar Hofmann
Closing Date	31 July, 2026

31 July, 2026

Chairman of the Examination Board

Contents

List of Abbreviations	iii
1 Introduction	1
2 Theoretical Background	3
2.1 Green AI and Red AI	3
2.2 Energy and Carbon Emissions in Machine Learning	4
2.2.1 From Energy to Carbon Emissions	4
2.2.2 Computing Hardware and Power Consumption	5
2.2.3 Measurement Approaches	6
2.2.4 CodeCarbon	7
2.3 Supervised Classification Algorithms	7
2.3.1 Logistic Regression	8
2.3.2 Decision Trees	9
2.3.3 Random Forest	10
2.3.4 Gradient Boosting and XGBoost	11
2.3.5 Multilayer Perceptron	14
2.3.6 Tree-based Models vs. Deep Learning on Tabular Data	15
2.3.7 Algorithmic Suitability for CPU vs. GPU	16
2.4 Dataset Scaling and Learning Curves	17
2.5 Evaluation Metrics	18
2.6 Cross-Validation	20
2.7 Hyperparameter Optimization	21
2.7.1 Grid Search and Random Search	22
2.7.2 Bayesian Optimization	22
2.7.3 Optuna as a Practical Framework	24
2.8 Chapter Summary	24
3 Related Work	27
3.1 Empirical Energy Benchmarks for Machine Learning	27
3.2 Dataset Size Effects on Performance and Energy Consumption	28
3.3 CPU vs. GPU Efficiency: Empirical Evidence	29
3.4 Measurement Tool Accuracy: Software Estimation vs. Hardware Sensors	30
3.5 Seasonal and Geographic Carbon Intensity Effects	31
3.6 Hyperparameter Tuning Efficiency	31
3.7 Lifecycle Emissions: The Ecological Cost of Deployment and Inference	32

3.8	Composite Efficiency Metrics: Balancing Performance and Ecological Cost . . .	33
4	Methodology	35
4.1	Research Questions	35
4.2	Datasets	35
4.3	Models	36
4.4	Evaluation Metrics	37
4.5	Validation Strategy	38
4.6	Hyperparameter Tuning	38
5	Implementation and Measurement	39
5.1	Experimental Setup	39
5.1.1	Hardware Setup	39
5.1.2	Software Environment and Reproducibility	40
5.2	Dataset Integration	41
5.3	Model Implementations	42
5.3.1	Logistic Regression	42
5.3.2	Random Forest	43
5.3.3	XGBoost	43
5.3.4	Multilayer Perceptron	44
5.4	Hyperparameter Optimization	45
5.5	Emissions Measurement	47
5.6	Inference Latency Measurement	49
5.7	Results Storage and Orchestration	50
5.8	Experiments	52
5.8.1	Main Cross-Dataset Benchmark	52
5.8.2	Dataset Scaling Experiment	52
5.8.3	MLP Architecture Variation	53
5.8.4	Temporal Carbon Intensity Analysis	54
5.8.5	Lifecycle Break-even Analysis	55
6	Results and Discussion	57
6.1	Emissions Measurement Validation	57
6.2	Hyperparameter Tuning Overhead	57
6.3	Main Cross-Dataset Benchmark	58
6.4	XGBoost: CPU vs. GPU	60
6.5	HIGGS Dataset Scaling	60
6.6	MLP Architecture Variation	62
6.7	Carbon Intensity Analysis	63
6.8	Lifecycle Break-Even Analysis	68
	Bibliography	69
	List of Figures	74
	List of Tables	75

List of Abbreviations

ACC Accuracy.

Adam Adaptive Moment Estimation.

AMP Automatic Mixed Precision.

CPU Central Processing Unit.

DRAM Dynamic Random-Access Memory.

EPM External Power Meter.

EWf1 Efficiency-Weighted F1.

FLOP Floating-Point Operation.

FP32 32-bit single-precision floating-point format.

GPU Graphics Processing Unit.

kgCO₂eq kilograms of CO₂-equivalents.

LOOCV Leave-One-Out Cross-Validation.

LR Logistic Regression.

MLP Multilayer Perceptron.

NN Neural Network.

NVML NVIDIA Management Library.

PMC Performance Monitoring Counter.

RAPL Running Average Power Limit.

RF Random Forest.

SGD Stochastic Gradient Descent.

TDP Thermal Design Power.

TN True Negative.

TP True Positive.

TPE Tree-structured Parzen Estimator.

WF1 Weighted F1-Score.

XGB Extreme Gradient Boosting.

Introduction

Theoretical Background

This chapter establishes the theoretical foundations required to understand the empirical study presented in the remainder of this thesis. It first frames the research within the broader debate between *Red AI* and *Green AI*, then formalizes the concepts of energy consumption and carbon emissions in machine learning. Subsequently, the relevant classification algorithms are described from a purely theoretical standpoint, followed by the metrics used to evaluate their predictive performance, the validation strategy, and hyperparameter optimization techniques. All implementation details are deliberately deferred to Chapter 5.

2.1 Green AI and Red AI

In recent years, a paradigm has emerged in AI research that Schwartz et al. refer to as "Red AI" [1]. Red AI describes research approaches that primarily aim to improve prediction accuracy through the massive use of computational power, data, and parameters, while largely disregarding the associated infrastructural and energetic costs. The underlying dynamic driving this trend is that the relationship between model performance and model complexity is at best logarithmic [1]: a linear increase in accuracy requires an exponentially larger model or exponentially more training data. Empirically, since 2012 the compute used to train state-of-the-art models has been doubling approximately every 3.4 months [2]. This is particularly evident in large Transformer architectures. Strubell et al. document that training such models can require days of continuous computation across multiple Graphics Processing Units (GPUs), while foundation models such as GPT-3 scale to 175 billion parameters [3].

In response, the concept of *Green AI* has emerged as a counter-movement within the machine learning community [1]. Green AI is defined as research that yields novel results while explicitly accounting for computational costs, elevating efficiency to a primary evaluation criterion alongside predictive accuracy. Beyond the environmental dimension, the movement also addresses inclusivity: the prohibitive financial costs of training massive models increasingly restrict participation to a small number of well-resourced institutions [1]. A distinction must also be made between Green AI and *AI for Green*, where Green AI concerns reducing the ecological footprint of AI systems themselves, not applying AI to environmental problems [1, 4].

A core principle of the Green AI movement is evaluating environmental impacts across the entire lifecycle of an AI system, rather than focusing solely on the training phase [4, 5]. This full-lifecycle view encompasses data engineering, system integration, continuous retraining, and inference operations. Empirical work on data-centric Green AI illustrates the potential of this broader perspective: simple dataset preprocessing modifications, such as

extracting smaller representative subsets or reducing feature counts, can decrease training energy consumption by up to 92 percent with negligible loss in accuracy [2].

To translate these principles into practice, researchers advocate for standardized reporting of computational work, energy consumption, and carbon emissions alongside accuracy results [1, 6]. Several tools have been developed to support this. The ML CO₂ Impact Calculator [7] and the experiment-impact-tracker [6] allow practitioners to estimate carbon emissions based on hardware configuration, execution time, and local grid carbon intensity. At the organizational level, the Green Software Foundation’s Software Carbon Intensity (SCI) specification provides a standardized, consequential carbon accounting metric aimed at directly reducing AI carbon footprints rather than relying on market-based offsets [5, 8].

Realizing efficiency as a measurable criterion requires defining appropriate metrics. Floating-Point Operations (FLOPs) provide a hardware-independent estimate of computational work [1], but do not translate directly to energy or execution time, as these depend on hardware architecture, memory access patterns, and parallelization [6]. Training time is more intuitive but hardware-dependent and difficult to compare across environments. Energy consumption in kWh and CO₂ equivalents (CO₂eq) sidestep these limitations by capturing the actual physical and ecological footprint of a training run. Although tied to a specific hardware configuration and electricity grid, they provide directly comparable sustainability figures when experiments are conducted under controlled conditions [6, 8].

2.2 Energy and Carbon Emissions in Machine Learning

Operationalizing efficiency requires distinguishing two related but distinct quantities. Energy consumption is a hardware-level physical quantity, the product of power draw and time expressed in kilowatt-hours, where the actual power draw depends on hardware utilization and thermal design rather than being a fixed property of the components [6]. Carbon emissions then translate that energy into an ecological footprint, measured in carbon dioxide equivalents (CO₂eq), by multiplying it with the *carbon intensity* of the local electricity grid [7]. Two identical training runs can therefore produce the same energy consumption yet vastly different emissions, depending solely on the fuel mix of the grid supplying them. In practice, empirical studies have shown that identical workloads can vary in their carbon footprint by a factor of 40 depending on geographic location [7], and can even fluctuate based on the time of day the models are trained [8].

2.2.1 From Energy to Carbon Emissions

To understand the ecological footprint of machine learning, it is first necessary to distinguish between power and energy. Power, measured in Watts (W) or Joules per second, represents the instantaneous rate at which electricity is drawn by a system. Energy, on the other hand, is the total quantity of electricity consumed over a specific duration, calculated as the product of power and time ($E = P \cdot t$), and is typically expressed in kilowatt-hours (kWh) or Joules [6].

Translating this measured energy consumption into a quantifiable environmental impact

requires the concept of carbon intensity [7]. Carbon intensity represents the amount of greenhouse gas emissions produced per unit of electricity generated, expressed in grams of CO₂ equivalents per kilowatt-hour (gCO₂eq/kWh), and standardizes the global warming potential of various greenhouse gases into a single comparable value.

The operational carbon footprint of a machine learning model is derived by multiplying the total energy consumed (E , in kWh) by the carbon intensity of the local electricity grid (I), yielding $\text{CO}_2\text{eq} = E \cdot I$ [8]. Because different geographic regions rely on vastly different fuel mixes, from coal and natural gas to nuclear and renewables, carbon intensity varies drastically by location [6]. Empirical data illustrates this range clearly: training a model in Quebec, Canada, which is heavily reliant on hydroelectricity, faces a carbon intensity of approximately 20 gCO₂eq/kWh, whereas training the same model in a fossil-fuel-dependent region such as Iowa, USA, results in a carbon intensity of over 736 gCO₂eq/kWh [7].

Carbon intensity is not merely a static geographic property but also a temporally variable one [8]. The emissions factor of a grid can fluctuate significantly throughout the day based on the real-time availability of renewable sources, such as solar during daylight hours or wind at night. Consequently, accurately evaluating the ecological footprint of a machine learning model requires integrating both the spatial location and the temporal characteristics of the local energy grid into the final CO₂eq conversion.

2.2.2 Computing Hardware and Power Consumption

The execution of machine learning workloads relies on a combination of hardware components, primarily the Central Processing Unit (CPU), GPU, and Dynamic Random-Access Memory (DRAM) [6]. Understanding the architectural differences between these components is essential for reasoning about their respective energy profiles and the efficiency trade-offs examined in this thesis.

The CPU is a general-purpose processor designed for sequential, low-latency execution of complex, branching instruction streams. Its architecture dedicates a large share of its transistor budget to control logic, branch prediction, and deep cache hierarchies, typically comprising between 4 and 32 high-frequency cores. This design excels at tasks requiring complex decision logic but offers limited parallelism for the arithmetic-intensive operations that dominate machine learning.

The GPU, by contrast, is a hardware accelerator originally designed for rendering graphics, a task that requires applying the same transformation to millions of pixels simultaneously. Its architecture instead consists of thousands of smaller, simpler processing cores organized into streaming multiprocessors, optimized for high arithmetic throughput and high memory bandwidth. This massively parallel design maps directly onto the matrix multiplications and element-wise operations that dominate both training and inference in machine learning, which is why GPUs have become the standard compute substrate for computationally intensive ML workloads [8].

This architectural advantage comes at a steep energetic cost. While CPUs typically draw between 10 W and 150 W under load, GPUs are significantly more power-hungry, often consuming between 250 W and 350 W [8]. As a result, the GPU typically accounts for

the vast majority of total electrical energy consumed during training, often approaching three-quarters of the total draw [8]. The energy drawn by the CPU and DRAM nevertheless remains a non-negligible contribution that must be recorded for accurate sustainability accounting.

A frequently referenced hardware specification in this context is the Thermal Design Power (TDP). TDP defines the maximum heat generated by a component that its cooling system is designed to dissipate, serving as a theoretical upper bound on sustained power draw [9]. Some simplified methodologies estimate total energy consumption by multiplying a component's TDP by execution time, but this approach is fundamentally flawed [6]: actual power draw fluctuates continuously based on hardware utilization, varying greatly between idle states and peak load. Empirical studies demonstrate that TDP-based estimation leads to severe over- or underestimations of the true energy footprint [6]. Reliable sustainability measurements therefore require direct real-time hardware readings rather than static TDP-based extrapolations.

2.2.3 Measurement Approaches

Accurately evaluating the energy consumption of machine learning models relies on a spectrum of methodologies, broadly categorized into physical measurements, estimation models, and on-chip sensors [9]. The baseline for ground-truth measurement is the External Power Meter (EPM), which tracks power directly at the wall outlet or motherboard. While EPMs capture the true consumption of the entire system, they are often costly to deploy at scale and cannot provide fine-grained decomposition to isolate the energy used by specific software processes or individual hardware components [9].

To achieve component-level granularity without external hardware, many approaches utilize estimation models. These models calculate energy based on abstract hardware specifications, such as multiplying execution time by the TDP or counting FLOPs, or by leveraging Performance Monitoring Counters (PMCs) to track utilization rates [9]. However, as established in Section 2.2.1, relying solely on abstract analytical models or TDP often leads to severe inaccuracies under dynamic workloads, failing to capture the true, fluctuating real-time power draw of the hardware [6].

Consequently, modern sustainability accounting strongly prefers direct measurements obtained from vendor-specific on-chip sensors [6]. For GPUs, the industry standard is the NVIDIA Management Library (NVML), which polls instantaneous power consumption during execution. For CPUs and DRAM, the standard interface is Running Average Power Limit (RAPL), provided by Intel and AMD, which accesses hardware performance counters to deliver accurate, fine-grained energy measurements directly from the CPU package [9].

Despite its high accuracy, a significant practical limitation of RAPL is its strict dependency on the underlying operating system [6]. RAPL interfaces are primarily supported on Linux environments and often require elevated administrator privileges to access the hardware counters securely. When direct sensor access is unavailable due to these constraints, software tracking tools often revert to problematic TDP-based estimation models as a fallback, which severely compromises measurement accuracy [9].

2.2.4 CodeCarbon

To practically implement the sustainability accounting principles advocated by the Green AI movement, researchers require standardized tools that bridge the gap between physical hardware metrics and environmental impacts. Initially, the machine learning community relied on a fragmented landscape of independent measurement approaches and preliminary software packages [10]. Early solutions largely consisted of web-based calculators, such as the ML CO₂ Impact calculator [7] and Green Algorithms [11], which required practitioners to manually input hardware specifications, execution times, and locations to estimate carbon footprints.

While these online calculators successfully raised awareness, they lacked real-time monitoring capabilities and had to rely on broad estimations rather than ground-truth physical power draw [10]. To address this, several independent Python packages were developed to track energy dynamically during model training by directly polling hardware sensors during execution, including the experiment-impact-tracker [6] and CarbonTracker [12]. However, the proliferation of tools with varying methodologies, carbon intensity data sources, and hardware support led to inconsistent reporting and discrepancies in carbon footprint measurements across the literature [10].

To unify these fragmented efforts into a reliable standard, the CodeCarbon initiative was established [10]. CodeCarbon is an open-source software library that evolved directly from previous tracking projects, serving as the official successor to the Energy Usage tool [9] and integrating the efforts of the ML CO₂ Impact project [10]. By providing a unified tracking interface, it enables practitioners to quantify the environmental cost of their algorithms and report these metrics transparently [13].

To evaluate a model's energy consumption, the framework leverages the direct measurement approaches introduced in Section 2.2.3. Specifically, it interfaces with vendor-specific on-chip sensors, such as NVML for GPUs and RAPL for Intel and AMD processors, to collect fine-grained power draw measurements directly from the underlying hardware [9]. By periodically polling these hardware counters during execution, the tool computes the total electrical energy consumed by the computing components in kilowatt-hours.

Once the physical energy consumption is recorded, the framework translates it into greenhouse gas emissions using the conversion detailed in Section 2.2. The tool determines the carbon intensity of the local electricity grid based on the geographic location of the computing infrastructure, then multiplies the measured energy by this value to produce the final operational carbon footprint expressed in kilograms of CO₂-equivalents (kgCO₂eq) [8]. By combining hardware-level energy tracking with grid-level carbon intensity data in a single standardized framework, CodeCarbon automates the translation from computational work to ecological impact [6].

2.3 Supervised Classification Algorithms

In machine learning, supervised classification is the task of inferring a predictive model from historical, annotated data in order to make accurate categorical predictions on new data

[14]. Formally, given a training dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ consisting of n examples, where each $x_i \in \mathcal{X}$ represents a vector of input features and $y_i \in \mathcal{Y}$ represents the corresponding ground-truth label, the objective is to learn a mapping function $f : \mathcal{X} \rightarrow \mathcal{Y}$ that reliably predicts the target output for any given input [15]. This reliance on an explicitly provided label space distinguishes supervised learning from unsupervised learning, where algorithms discover hidden structures in unannotated data, and from reinforcement learning, where an agent learns to make sequential decisions by maximizing a reward signal within an interactive environment.

The fundamental objective when learning f is achieving strong *generalization*: the model must perform reliably on new, unseen data drawn from the same underlying distribution rather than merely memorizing training inputs [16]. A central challenge during training is therefore finding the optimal balance to prevent *underfitting*, where the model is too simple to capture the true data patterns, and *overfitting*. Overfitting occurs when a model excessively adapts to the specific “sampling noise” present in the training data, fitting it so closely that its predictive performance on new, unseen data actively degrades [16, 17].

The structure of the label space $\mathcal{Y}$ determines the nature of the classification task. In binary classification, $\mathcal{Y} = \{0, 1\}$. When more than two categories are involved it becomes a multiclass problem with $\mathcal{Y} = \{1, \dots, C\}$ [15]. Multiclass problems are typically handled via One-vs-Rest strategies or by extending binary formulations with a Softmax output that produces a normalized probability distribution across all C classes.

To optimize f during training, a loss function $\mathcal{L}(y, f(x))$ quantifies the discrepancy between predicted outputs and true labels [17]. For classification, the standard objective is Cross-Entropy Loss. In the binary case this reduces to the negative binomial log-likelihood, while for multiclass problems it generalizes to Categorical Cross-Entropy, formalized in Equation (2.1).

$$\mathcal{L} = - \sum_{k=1}^C y_k \log(p_k(x)) \quad (2.1)$$

In Equation (2.1), y_k is the binary indicator for the true class k and $p_k(x)$ is the model’s predicted probability for that class. Cross-Entropy is preferred over Mean Squared Error for classification because it produces steeper, more informative gradients when a confident prediction is wrong, and provides a direct probabilistic interpretation of the model output. This foundational loss objective drives the gradient-based optimization for both the deep learning and advanced ensemble architectures discussed in the following subsections.

2.3.1 Logistic Regression

Logistic regression, formalized as a statistical classification method by Cox [18], is a parametric linear model that remains a standard baseline for binary classification tasks where interpretability and computational efficiency are paramount.

For a binary classification task ($\mathcal{Y} = \{0, 1\}$), logistic regression predicts the probability that an instance x belongs to the positive class. To ensure the output is bounded strictly between 0 and 1, the linear combination of input features and learned weights w with bias b is passed through the logistic sigmoid function $\sigma(z) = \frac{1}{1+e^{-z}}$, yielding the prediction model

in Equation (2.2).

$$P(y = 1 | x) = \sigma(w^\top x + b) = \frac{1}{1 + e^{-(w^\top x + b)}} \quad (2.2)$$

This functional form stems from modeling the log-odds (logit) as a linear function of the inputs, as expressed in Equation (2.3) [14].

$$\ln\left(\frac{P(y = 1 | x)}{P(y = 0 | x)}\right) = w^\top x + b \quad (2.3)$$

Parameters are estimated via Maximum Likelihood Estimation, which is equivalent to minimising the binary cross-entropy loss and requires iterative numerical solvers since no closed-form solution exists [14]. This study uses the Stochastic Average Gradient (SAG) solver [19], which is particularly well-suited for large datasets: rather than recomputing the full gradient over all training samples at each step, SAG maintains a running average of recently computed individual gradients, substantially reducing both memory usage and required training time. For multiclass problems, Multinomial Logistic Regression with a Softmax function is applied [14].

A key advantage of logistic regression is its interpretability: each feature weight directly represents its directional influence on the log-odds of the prediction. It also trains quickly and is robust to overfitting on simple datasets. Its fundamental limitation is linearity: without manual feature engineering, the model cannot capture non-linear patterns or feature interactions, making it considerably less expressive than the tree-based and deep learning architectures described in the following subsections [14].

2.3.2 Decision Trees

Decision trees represent a fundamental family of non-parametric machine learning algorithms that approach classification through a recursive binary partitioning strategy [20]. In this architecture, learning proceeds top-down from a root node through internal nodes, each of which applies a threshold test of the form $x_j \leq t$ on a single input feature j with a learned constant threshold t . This threshold determines which branch a given data point follows: samples satisfying the condition are routed left, all others right. The tree routes each sample through a sequence of such decisions until it reaches a terminal leaf node. In essence, the tree divides the input feature space into distinct regions, with each region corresponding to a leaf. All data points that reach the same leaf receive the same class prediction: the majority class among the training samples that were routed there [17].

The induction of a decision tree relies on a greedy split-finding algorithm, which makes locally optimal decisions at each node without revisiting or reconsidering earlier splits. At each internal node, the algorithm evaluates candidate split points across all available features to find the partition that minimizes impurity in the resulting child nodes. The Gini impurity criterion is used, defined as $G = \sum_{k=1}^C \hat{p}_k(1 - \hat{p}_k)$, which measures the probability of misclassifying a randomly drawn sample [17]. This is the default criterion in scikit-learn [19].

The structural design of decision trees imparts specific inductive biases that make them particularly well-suited for tabular data, as documented by Grinsztajn et al. [21]. Because

they naturally learn piecewise constant functions, decision trees can fit highly irregular, non-smooth patterns without requiring complex mathematical transformations. Furthermore, decision trees operate strictly through axis-aligned splits and are therefore not rotationally invariant. This is highly advantageous for tabular data, as it prevents the algorithm from creating arbitrary linear combinations of variables, thereby preserving the independent, heterogeneous meaning of the original features [21]. Consequently, trees remain robust even when datasets contain uninformative features [22].

Despite these representational strengths, single decision trees suffer from significant theoretical limitations. When grown deeply, a tree becomes highly susceptible to memorizing sampling noise, leading to high variance, where predictions fluctuate unpredictably across different training sets and severe overfitting, where the model fits training noise rather than underlying patterns, causing poor performance on new data [17]. Structural constraints such as a maximum depth or minimum number of samples per leaf partially mitigate this, but at the cost of increased bias, where the model becomes too restrictive and cannot capture underlying patterns and the need for careful tuning. To more systematically address these vulnerabilities while preserving the tree's ability to model irregular tabular data, the machine learning community shifted toward combining multiple tree predictors into ensemble methods [22]. This transition from single, unstable base learners to aggregated models forms the foundational motivation for the Random Forest and Gradient Boosting architectures described in the following sections.

2.3.3 Random Forest

As established in the previous section, the inherent instability, high variance, and overfitting risk of single decision trees serve as the primary motivation for aggregating multiple individual trees into powerful ensemble methods [22]. Random Forests, formally introduced by Leo Breiman [22], overcome these limitations by combining a large collection of tree predictors, where each tree depends on the values of a random vector sampled independently and with the same distribution for all trees in the forest.

The algorithm constructs this ensemble using a two-fold randomization strategy. First, it employs a technique known as *bagging* (bootstrap aggregating), originally proposed by Breiman [23], where each tree is grown on a new training set drawn at random with replacement from the original data. Random Forests extend this by additionally randomizing feature selection during the tree-growing phase: instead of evaluating all available features to determine the optimal split at each internal node, the algorithm selects the best split from a randomly chosen, restricted subset of features [22]. The constituent trees are typically grown to their maximum depth and are not pruned. To classify a new input vector, each tree in the ensemble casts a unit vote, and the forest outputs the majority class.

Breiman theoretically proved that the predictive accuracy of a Random Forest is governed by the interplay of two key parameters: the *strength* of the individual tree classifiers and the *correlation* between them [22]. Correlation here means the degree to which trees make similar predictions on the same data points. Decorrelated trees, by contrast, make diverse predictions and thus commit different kinds of errors. An optimal forest consists of

individually strong trees that produce decorrelated predictions. By combining bagging with random feature selection, the algorithm deliberately decreases the correlation between individual trees, which substantially reduces the variance of the overall ensemble without substantially compromising the predictive strength of the base learners. Building on this analysis, Breiman further demonstrated via the Strong Law of Large Numbers that as the number of trees grows, the generalization error almost surely converges to a limiting value, guaranteeing that the ensemble is inherently robust against overfitting with respect to the number of trees. While this convergence argument in the original paper was somewhat informal, formal consistency for a close variant of Breiman's algorithm under standard assumptions was later established by Scornet, Biau, and Vert [24].

A crucial computational advantage of this architecture, particularly relevant for evaluating the efficiency of machine learning algorithms, is its inherent parallelizability [25]. Because the individual trees are entirely independent of one another, the training process can be distributed independently across multiple processing cores, allowing Random Forests to scale efficiently on multi-core CPU architectures [22, 25]. GPU-accelerated implementations also exist (e.g., RAPIDS cuML, Intel Extension for Scikit-learn), though GPU support remains implementation-dependent rather than a universal property of the algorithm.¹ Additionally, the bagging mechanism naturally provides an internal validation mechanism: since each tree is trained on a bootstrap sample (approximately 63% of the data), the remaining samples—termed *out-of-bag* (OOB) samples—form the complement set. For each training sample, the forest can evaluate it using only the subset of trees that did not include it in their training bootstrap, yielding an unbiased generalization error estimate without requiring a separate validation split [22]. In scikit-learn, this mechanism is optional and requires setting the `oob_score=True` parameter at initialization.

Despite this parallelism, Random Forests carry a significant computational burden on large datasets. Each tree must be grown independently on a bootstrap sample of the full training set, and because the trees are deep and unpruned, meaning they are grown to full depth without being simplified or trimmed back, the cost per tree grows with both the number of samples and the number of features evaluated at each split [25]. This cost is then multiplied across all M trees in the ensemble. As the dataset size grows into the millions of samples, the aggregate training cost becomes substantial, limiting the practical applicability of Random Forests on very large datasets without subsampling or other approximations.

2.3.4 Gradient Boosting and XGBoost

Random Forests and Gradient Boosting represent fundamentally different ensemble strategies. Random Forests combine many independent strong learners: deep, complex trees trained in parallel on bootstrap samples, where the ensemble strength comes from diversity and aggregation. In contrast, Gradient Boosting operates sequentially, constructing an ensemble of weak learners (shallow trees) where each successive tree is trained to correct the errors of the previous iterations, iteratively improving overall performance [17]. Introduced

¹GPU acceleration for Random Forests is not universally available across all libraries. Support depends on the specific implementation: scikit-learn runs on CPUs only, while specialized libraries like RAPIDS cuML provide GPU-accelerated versions.

by Friedman [17], the gradient boosting machine frames this sequential training as a form of gradient descent in function space. Instead of adjusting numerical parameters, the algorithm iteratively adds a new decision tree to the ensemble that is fitted to the negative gradient of the loss function with respect to the predictions of the previous iteration. This allows the model to minimize arbitrary differentiable loss functions, making it a highly flexible and powerful approach for both regression and classification tasks. Throughout this framework, $\hat{y}_i$ denotes the model's predicted output for training instance i , where the hat symbol ($\hat{\cdot}$) is standard notation distinguishing a model's prediction from the corresponding true label y_i . After each boosting round, a new tree is added whose outputs correct the residual errors of the current ensemble, progressively updating $\hat{y}_i$ toward y_i . Figure 2.1 illustrates the architectural contrast between the two strategies: the left panel shows Random Forest's parallel structure, in which all trees are trained independently on bootstrap samples and their predictions are aggregated by majority vote, while the right panel shows Gradient Boosting's sequential pipeline, in which each successive tree is fitted to the errors of the previous ensemble and the predictions are summed to progressively reduce the residual.




Figure 2.1: Random Forest vs Gradient Boosting: Parallel (breadth) vs Sequential (depth) ensemble strategies.

XGBoost (Extreme Gradient Boosting), introduced by Chen and Guestrin [26], significantly extends this foundational concept to improve both predictive accuracy and computational scalability. To combat the overfitting risks inherent in boosted ensembles (where each tree iteratively fits residual errors of its predecessors, risking the memorization of training

noise), XGBoost modifies the training objective by augmenting the task-specific loss function with an explicit regularization term [26].

$$\mathcal{L} = \sum_{i=1}^n l(y_i, \hat{y}_i) + \Omega(f), \quad \Omega(f) = \gamma T + \frac{1}{2} \lambda \|w\|^2 \quad (2.4)$$

In Equation (2.4), n is the number of training instances, y_i the true label and $\hat{y}_i$ the predicted value for instance i , and $l(\cdot)$ is a differentiable task-specific loss function. For binary classification, for example, l takes the form of the logistic loss $l(y_i, \hat{y}_i) = y_i \log \hat{y}_i + (1 - y_i) \log(1 - \hat{y}_i)$. The regularization term $\Omega(f)$ penalizes model complexity directly: T is the number of leaf nodes of tree f , $w \in \mathbb{R}^T$ the vector of leaf output weights, $\gamma \geq 0$ a penalty controlling tree size, and $\lambda \geq 0$ the L2 regularization coefficient that shrinks the leaf weights toward zero to reduce variance. Furthermore, while standard gradient boosting relies only on first-order gradients, XGBoost utilizes a second-order Taylor expansion of the loss function [26]. By calculating both the gradient and the Hessian (second derivative) for each training instance, the algorithm can optimize the tree structure more directly and rapidly converge on the optimal split points.

XGBoost also implements *shrinkage* and *column subsampling* to further prevent overfitting [26]. Shrinkage scales each new tree's contribution by a learning rate to slow convergence, which is enabled by default in the XGBoost package via `eta=0.3`. Column subsampling introduces feature randomization analogous to Random Forest, but whereas Random Forest draws a fresh random feature subset at every individual split node, XGBoost's `colsample_bytree` selects a random subset of features once per tree, so that all nodes within that tree compete only over the same fixed subset, though this mechanism must be explicitly configured, as the default value of `colsample_bytree=1.0` retains all features and effectively disables it.

Despite these algorithmic optimizations, the exact greedy split-finding procedure shared by all tree-based methods requires sorting continuous feature values across all candidates, a bottleneck that becomes prohibitive at the scale of millions of instances [26]. To accelerate this, XGBoost was extended to natively leverage Graphics Processing Units (GPUs) [27]. Formalized by Mitchell and Frank, the GPU-accelerated tree construction algorithm exploits the massive parallelism and high memory bandwidth of modern GPU architectures. Instead of relying on exact sorting, the *histogram method* bins continuous feature values into discrete buckets and aggregates gradient statistics per bin. By evaluating split points using highly optimized parallel primitives over these histograms, the entire decision tree can be constructed directly in device memory [27]. This histogram-based split-finding approach allows XGBoost to scale efficiently to millions of instances, achieving significant speedups over multi-core CPU implementations while fully accommodating the algorithm's second-order mathematical formulations. In their evaluation, Mitchell and Frank demonstrated the memory efficiency of this approach by processing a large-scale dataset of 11 million training instances and 28 features entirely within the memory of a single GPU, confirming its practical viability for large-scale classification tasks [27]. This speedup, however, does not constitute a genuine efficiency gain in the Green AI sense: as Schwartz et al. [1] argue,

hardware acceleration merely redistributes computational cost from time to power draw, and true efficiency must come from algorithmic improvements that reduce the total work performed rather than from procuring faster hardware [1].

2.3.5 Multilayer Perceptron

The Multilayer Perceptron (MLP) is a feedforward neural network architecture composed of sequential layers of neurons that apply linear transformations followed by non-linear activation functions. The Universal Approximation Theorem [28] provides the mathematical justification for this architecture, guaranteeing that a single hidden layer is sufficient to approximate any continuous function arbitrarily closely, given enough neurons in that layer. However, unlocking this immense representational capacity requires scaling the network's depth and width, which inherently increases both the model's susceptibility to overfitting and the computational burden of training.

Unlike decision trees, MLPs are optimized continuously using *backpropagation* [29], which propagates the gradient of the loss function backwards through the network layer by layer via the chain rule to compute each parameter's contribution to the error, and Stochastic Gradient Descent (SGD) or its variants to minimize a predefined loss function. In practice, SGD is applied in its mini-batch form, where the gradient is computed not over the full dataset nor a single sample, but over small random subsets of the training data, trading gradient accuracy for computational efficiency and introducing beneficial noise that helps escape local minima. To accelerate this optimization process in modern applications, the Adaptive Moment Estimation (Adam) optimizer is frequently deployed. Adam is a computationally efficient, first-order gradient-based optimization algorithm [30]. It computes individual adaptive learning rates for different parameters from estimates of the first and second moments of the gradients. By maintaining exponential moving averages of the gradient and the squared gradient, and applying an initialization bias correction to counteract initial estimates being biased toward zero, Adam effectively adapts to the geometry of the objective function.

Historically, training deep architectures was severely hindered by the *vanishing gradient problem*. When utilizing traditional saturating non-linearities like the logistic sigmoid or hyperbolic tangent, the gradients of saturated neurons rapidly approach zero during backpropagation, making it nearly impossible for earlier layers to update effectively [31]. The widespread adoption of the Rectified Linear Unit (ReLU) largely resolved this bottleneck. By computing $y = \max(0, x)$, ReLU prevents gradients from vanishing for positive inputs and yields highly sparse representations [32]. Furthermore, its gradient is exactly 1 for positive inputs and 0 otherwise, making it highly efficient for training deep networks.

As MLPs scale in depth, they introduce considerable internal instability during training. During gradient descent, the distribution of inputs to any given hidden layer constantly changes as the parameters of all preceding layers are updated, a phenomenon termed *internal covariate shift* [33]. This forces the optimizer to use excessively low learning rates, substantially slowing down convergence. Batch Normalization overcomes this by explicitly standardizing the intermediate activations across each mini-batch to maintain a stable

mean of zero and variance of one, before applying a learned affine transformation with parameters γ and β [33]. By fixing these distributions, Batch Normalization allows for the use of substantially higher learning rates and makes backpropagation highly resilient to parameter scale, resulting in substantially faster training times.

To prevent these deep networks from merely memorizing the sampling noise of the training data, regularization is required. *Dropout* provides a computationally elegant solution by randomly deactivating a fraction of hidden units during each forward pass [16]. This prevents neurons from developing complex co-adaptations where they solely rely on specific other units to correct their mistakes. At test time, scaling the outgoing weights seamlessly approximates an equally weighted geometric mean of all these shared-weight subnetworks, significantly reducing generalization error without the prohibitive cost of training multiple large models [16].

Similarly, because an MLP will eventually begin to overfit if training continues without constraint, *Early Stopping* is employed to monitor a separate validation set and halt the training process once the validation error stops decreasing [34]. Beyond acting as an implicit regularizer, this technique serves as a critical efficiency mechanism, saving significant computational time by avoiding redundant training epochs.

A significant hardware-aware optimization for training deep neural networks is Automatic Mixed Precision (AMP). Standard training typically relies on the 32-bit single-precision floating-point format (FP32) for all calculations. AMP dynamically transitions between FP32 and 16-bit half-precision (FP16 or bfloat16) formats during execution. Operations that are highly sensitive to numerical underflow, such as gradient accumulation or specific loss computations, are preserved in FP32 to maintain numerical stability. Simultaneously, computationally dense operations, such as matrix multiplications, are cast to lower precision. This selective reduction in precision effectively halves the memory bandwidth required for those operations and allows the model to leverage specialized GPU hardware, such as Tensor Cores, which are explicitly optimized for low-precision arithmetic. Consequently, AMP can substantially accelerate training times and reduce the overall energy footprint of deep learning architectures without causing a perceptible degradation in the final predictive accuracy.

2.3.6 Tree-based Models vs. Deep Learning on Tabular Data

While deep learning has enabled tremendous progress on text and image datasets, its superiority on tabular data remains unproven [21]. For problems with data stored in tabular formats consisting of heterogeneous features, ensembles of decision trees, particularly Gradient Boosted Decision Trees (GBDT) and Random Forests, are consistently recognized as the leading tools for both practitioners and machine learning competitions [15]. Extensive recent benchmarking has confirmed that tree-based models remain the superior choice on medium-sized tabular datasets, frequently outperforming deep learning architectures even before accounting for their superior training speed. However, this performance gap narrows significantly as the dataset size increases. Tree ensembles dominate in low-data regimes, but deep learning models become increasingly competitive when the number of training

instances scales into the millions [21].

The performance gap between neural networks and tree-based models stems from their fundamentally different inductive biases and data requirements [21]. Neural networks are intrinsically biased toward overly smooth, low-frequency solutions. Consequently, they struggle to learn the highly irregular and non-smooth target functions that frequently characterize tabular datasets. In contrast, models based on decision trees learn piecewise constant functions and do not exhibit this smoothing bias, making them naturally adept at fitting these irregular patterns [21]. Beyond these structural differences, neural networks also exhibit poor sample efficiency compared to tree-based ensembles: deep learning models generally require substantially more data to learn useful representations without overfitting, making them ill-suited for very small tabular datasets where algorithms like Random Forests can quickly capture the underlying structure [21].

Furthermore, tabular datasets typically contain a large proportion of uninformative features [21]. Standard deep learning architectures, such as MLP, are highly susceptible to these uninformative features, whereas tree-based models remain robust even when up to half of the dataset's features carry no information. Another critical distinction is rotation invariance. Neural networks are typically rotationally invariant, which degrades performance on tabular data because individual tabular features generally carry specific, distinct meanings. Applying rotational transformations mixes features with very different statistical properties, destroying the natural orientation of the data [21]. Conversely, the axis-aligned splits of decision trees preserve this information by avoiding arbitrary linear combinations of distinct features [21].

Despite these challenges, there has been significant research into designing deep learning architectures specifically tailored for tabular data [15]. Examples include differentiable tree ensembles and adaptations of the Transformer architecture, such as FT-Transformer, which applies self-attention mechanisms directly to feature embeddings. While these advanced models close the performance gap and achieve state-of-the-art results among deep networks, they still do not universally outperform GBDT implementations like XGBoost across all tasks, especially when hyperparameters are properly tuned [15]. Crucially, even when complex neural networks manage to match or slightly exceed the predictive accuracy of tree-based ensembles on datasets of millions of instances, they demand substantially more computational resources and training time to achieve these results [21]. This fundamental trade-off between predictive gain and computational cost reinforces the need to evaluate models not solely on accuracy, but also on their overall resource and energy efficiency.

2.3.7 Algorithmic Suitability for CPU vs. GPU

In the context of Green AI, optimizing energy efficiency requires aligning the computational profile of a machine learning algorithm with the architectural strengths of the underlying hardware [1]. While it is common practice to universally accelerate training using GPUs, this approach is not intrinsically energy-efficient for all algorithms. The fundamental divide in hardware suitability lies in the distinction between dense, regular parallelism and irregular, sequential branching.

Hardware platforms are designed with opposing optimization goals: GPUs are optimized for aggregate throughput on parallel tasks, whereas CPUs are optimized for low latency on sequential, single-threaded tasks [35]. GPUs achieve this through a *Single Instruction, Multiple Threads* (SIMT) architecture that executes thousands of lightweight threads concurrently to hide memory latency. This design heavily favors dense, regular computational tasks with uniform memory access patterns, most notably the large-scale matrix multiplications that form the computational core of neural network training. When operations can be uniformly distributed across thousands of cores simultaneously, the GPU achieves maximum utilization.

Conversely, CPUs rely on large caches, branch prediction, and speculative execution to handle irregular control flows and sequential decision logic with minimal latency. Algorithms that depend on dynamic branching, non-contiguous memory access, or sequential data traversal perform poorly on GPUs: in a SIMT environment, divergent threads within a warp must serialize, negating the throughput advantage and leaving cores idle [27]. For these workloads, the CPU is the more efficient hardware choice.

Despite the higher peak power draw of GPUs, their throughput on suitable workloads means they process substantially more training samples per joule than a CPU running the same task. A GPU can therefore deliver superior energy efficiency for well-matched workloads despite consuming more power in absolute terms. This advantage is not universal: at small dataset scales, the GPU's higher idle and transfer overhead can make CPU training the more energy-efficient option overall.

2.4 Dataset Scaling and Learning Curves

In predictive modeling, the relationship between the volume of available training data and the resulting classification performance is formally modeled using a learning curve [36]. Empirical learning curves consistently demonstrate that performance improves as more data becomes available, but the rate of improvement eventually diminishes, typically following an inverse power law $P(n) \propto n^\alpha$ with $\alpha \in (0, 1)$ [14, 36]. This creates a fundamental tension in the context of Green AI: while computational and energy costs scale roughly linearly with dataset size, performance yields only sub-linear growth, making it increasingly ecologically expensive to achieve stronger results simply by adding more data [1]. Achieving a marginal accuracy gain at the upper end of a learning curve demands a disproportionately large investment in additional data, time, and energy.

Conversely, when dataset sizes are heavily restricted, classifiers generally experience degraded performance because smaller datasets contain fewer representative examples, leaving the model unable to generalize underlying patterns [37]. The risk of overfitting also increases: with limited samples, a model may memorize training noise rather than learn true structure. Crucially, however, the overall performance of a classifier depends more heavily on how well the dataset represents the underlying data distribution than on its sheer size [37]. A small but highly representative dataset can yield better generalization than a large but biased one.

Different classes of machine learning algorithms exhibit different sensitivities to dataset

scale. Deep learning architectures are highly data-dependent: they must fit a large number of parameters via backpropagation, meaning more data directly enables better adjustment and generalization [37]. Tree-based models are also sensitive at the smaller end of the scale, as recursive partitioning means predictions at deeper nodes rest on increasingly small subsets of examples, making tree depth unstable when the initial dataset is limited [22, 37]. In contrast, simpler linear models serve as robust baselines: their constrained functional form imparts a higher model bias but a much lower variance, making them far less prone to overfitting on limited datasets [14].

A complex model’s absolute accuracy must therefore be weighed against its computational footprint and the diminishing returns of the learning curve. This trade-off is especially pronounced in a Green AI context, where the ecological cost of scaling up model complexity or dataset size may exceed the marginal accuracy gain.

2.5 Evaluation Metrics

Assessing the ecological efficiency of a machine learning algorithm requires more than measuring its predictive performance in isolation. An algorithm that achieves high accuracy at disproportionate computational cost is not unconditionally preferable to one that is slightly less accurate but significantly more resource-efficient [1]. This section therefore introduces the theoretical basis for three families of metrics. The first, classification quality, measures how accurately a model predicts class labels. The second, resource consumption, quantifies the energy expenditure and associated carbon emissions of training. The third, composite efficiency, integrates both dimensions into a single score that captures the trade-off between predictive performance and ecological cost. The specific metrics and their experimental implementation are described in Section 4.4.

Confusion Matrix. At the core of evaluating classification quality is the *confusion matrix*, which cross-tabulates a model’s predictions against the true labels. For a binary classification problem, this matrix consists of four fundamental components: True Positives (TP), True Negatives (TN), False Positives (FP), and False Negatives (FN). This concept extends naturally to multiclass problems by tracking predictions across a $C \times C$ grid, forming the basis for deriving all subsequent performance metrics [14].

Table 2.1: Binary confusion matrix with the four fundamental prediction outcomes.

	Predicted Positive	Predicted Negative
Actual Positive	True Positive (TP)	False Negative (FN)
Actual Negative	False Positive (FP)	True Negative (TN)

Accuracy. The most intuitive of these metrics is accuracy, defined as the overall proportion of correct predictions: $\text{Accuracy} = \frac{TP+TN}{N}$, where True Positive (TP) and True Negative (TN) are the correctly classified instances of each class, and N is the total number of instances. While simple to interpret, accuracy exhibits a critical vulnerability when applied to datasets

characterized by an uneven distribution of labels, known as *class imbalance*. For example, in a dataset where 99% of the instances belong to a single majority class, a trivial classifier that ignores all input features and simply predicts the majority class every time will achieve 99% accuracy. Despite this high score, the model learns no underlying patterns and fails completely at identifying the minority class.

Precision, Recall, and F1-Score. This limitation motivates the use of more robust metrics, specifically precision and recall. Precision ($\frac{TP}{TP+FP}$) measures the proportion of positive predictions that are actually correct, while recall ($\frac{TP}{TP+FN}$) measures the proportion of actual positive instances successfully identified. Unlike accuracy, precision and recall evaluate performance entirely independently of true negatives, making them highly resilient to datasets where a large, heterogeneous majority class dominates the negative counts [38]. Because improving precision often comes at the expense of recall and vice versa, the F1-score combines them into a single metric calculated as their harmonic mean, formalized in Equation (2.5).

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (2.5)$$

The harmonic mean is explicitly chosen over the arithmetic mean because it heavily penalizes extreme values. A model that achieves near-perfect recall but abysmal precision will yield a low F1-score, forcing the classifier to balance both dimensions effectively. Consequently, the F1-score provides a more encompassing summary of overall model performance and overcomes the representation problems caused by imbalanced data distributions.

F1 Aggregation Strategies. When dealing with multiclass or imbalanced datasets, the F1-score can be aggregated in several ways. Macro F1 calculates the metric independently for each class and averages them, treating all classes equally regardless of their size [38]. Micro F1 aggregates the global contributions of TP, FP, and FN across all classes before calculating the score, an approach that inherently favors larger classes [38]. Weighted F1 calculates the score for each class independently but averages them using the true support of each class as a weight.

Among these aggregation strategies, Weighted F1 is particularly well-suited to imbalanced settings. By weighting each class's score according to its true support, it yields a more representative picture of overall model performance than either Macro F1, which gives disproportionate influence to rare classes, or simple accuracy, which is entirely dominated by the majority class.

Resource Metrics. Measuring classification quality, however, captures only one dimension of model evaluation. Evaluating models under the Green AI paradigm requires moving beyond pure predictive quality to measure the physical and environmental costs of achieving that performance [1]. The primary resource metrics tracked include training time (measured in seconds), energy consumption (measured in kilowatt-hours, kWh), operational carbon emissions (measured in kg of CO₂eq), and inference latency (seconds per sample). The foundational definitions and physical measurement approaches for these quantities are

established in Section 2.2.

Composite Efficiency Metrics. To evaluate machine learning models across both dimensions at once, researchers increasingly employ composite efficiency metrics that weigh predictive performance against these resource expenditures in a single score. The general methodological motivation behind these composite metrics is to provide a unified heuristic for identifying algorithms that offer the optimal trade-off between classification quality and ecological impact [1]. In the literature, such approaches include the Energy-Efficiency Ratio, Performance-per-Watt [39], and FLOPs-per-Sample (or per-inference) [1].

However, composite metrics possess inherent theoretical weaknesses. First, the resource-based components of these metrics, such as energy consumed or execution time, are intrinsically hardware-dependent. A composite score calculated in one environment is tied to the specific computational architecture and local energy grid, meaning it cannot be universally generalized or directly compared across different settings [1, 8]. Second, the weighting between performance and resource cost is highly scenario-dependent. Determining how much a marginal gain in predictive accuracy is worth compared to the energy required to achieve it remains an open question that depends heavily on the specific priorities or economic utility of the deployment scenario [6].

2.6 Cross-Validation

The simplest approach to evaluating a model's generalization capability is the *hold-out method*, where the available data is partitioned into a single training and test split (for example, an 80/20 ratio). While computationally efficient, this approach is highly sensitive to the specific data partition and exhibits high variance in its performance estimate, particularly when the dataset is small [14]. To mitigate this instability, *K-fold cross-validation* is widely employed as a more rigorous evaluation technique [14].

The algorithm partitions the dataset into K roughly equal-sized folds. The model is trained K times. In each iteration, $K - 1$ folds are used to fit the model, while the single held-out fold is used for testing [14]. The cross-validation estimator is formalized in Equation (2.6).

Im Meeting
besprochen

$$CV(K) = \frac{1}{K} \sum_{k=1}^K L(y_k, \hat{f}^{-k}(x_k)) \quad (2.6)$$

In Equation (2.6), L is the chosen evaluation metric, y_k are the true labels of fold k , x_k are the corresponding input features of fold k , and $\hat{f}^{-k}$ denotes the model trained on all folds except the k -th [14]. The standard deviation of L across all K folds provides a quantitative measure of the model's stability across different data samples.

Selecting the number of folds K introduces a fundamental bias-variance trade-off. A small value of K is computationally less expensive but yields a higher estimation bias, as the model is trained on a significantly smaller fraction of the original data [14]. Conversely, setting $K = N$, known as Leave-One-Out Cross-Validation (LOOCV), trains the model N times using $N - 1$ samples for each iteration. While LOOCV produces a nearly unbiased

estimate of the test error, it is computationally prohibitive for large datasets and can exhibit high variance. Because the N training sets differ by only a single sample, the resulting estimators are highly correlated. Unlike independent estimators, whose averaged variance shrinks by a factor of $\frac{1}{N}$, the variance of correlated estimates does not diminish at this rate, so the averaging step of LOOCV provides little variance reduction [14]. To balance these competing factors, $K = 5$ and $K = 10$ are established as the standard recommendations in statistical learning, offering an optimal compromise between computational feasibility, low bias, and acceptable variance [14]. The specific justification for $K = 5$ is detailed in Section 4.5.

Finally, when evaluating datasets characterized by class imbalance, standard random partitioning may result in folds that do not accurately represent the underlying target distribution. To address this, *stratified cross-validation*, a variant that ensures the original class ratio is preserved within every fold, is applied. This is directly consistent with the choice of Weighted F1 as the primary evaluation metric: the same class imbalance that motivates a support-weighted metric must also be faithfully reproduced in each fold's test set, ensuring the evaluation procedure and the metric it produces remain coherent.

2.7 Hyperparameter Optimization

Machine learning models are governed by two distinct categories of configurable quantities. Parameters, such as the weights of a neural network or the split thresholds of a decision tree, are internal to the model and are learned directly from training data by minimizing a loss function [40]. Hyperparameters, by contrast, are external configuration choices that must be fixed prior to training: examples include the learning rate of an optimizer, the maximum depth of a tree, or the number of estimators in an ensemble [40]. Because hyperparameters govern the model's capacity, its regularization strength, which controls the degree to which overfitting is penalized, and its optimization dynamics, which determine how the learning process converges toward a solution, their values have a direct and often decisive impact on generalization performance [40, 41].

The goal of hyperparameter optimization (HPO) is to identify the configuration from a predefined search space Λ that minimizes the expected generalization error, where Λ is the Cartesian product of the individual value ranges of each hyperparameter, as shown in Equation (2.7), and the error is estimated via cross-validation [41]

$$\Lambda = \Lambda_1 \times \Lambda_2 \times \cdots \times \Lambda_d \quad (2.7)$$

where Λ_i denotes the value range of the i -th hyperparameter and d is the total number of hyperparameters. As the number of hyperparameters grows, exhaustive evaluation of Λ becomes exponentially expensive [40, 41]. This combinatorial explosion motivates the development of more efficient search strategies, which are described in the following subsections.

2.7.1 Grid Search and Random Search

The most traditional approach to hyperparameter tuning is Grid Search, an exhaustive strategy that evaluates the Cartesian product of a user-specified, finite set of values for each parameter [40]. While conceptually simple, Grid Search suffers severely from the *curse of dimensionality* [40, 41]. As the number of hyperparameters increases, the computational complexity grows exponentially. For instance, optimizing five hyperparameters with ten candidate values each across four models and three datasets would require an exhaustive evaluation of 150,000 distinct configurations. This combinatorial explosion renders Grid Search computationally prohibitive for complex machine learning architectures with limited budgets [40].

To overcome the inefficiency of exhaustive evaluation, Random Search was proposed as a highly effective alternative [41]. Instead of testing predetermined values on a rigid grid, Random Search selects a predefined number of parameter combinations randomly from the specified search space distributions [40]. Empirical and theoretical analyses demonstrate that Random Search is significantly more efficient than Grid Search in high-dimensional spaces [41]. The key insight is that the mapping from hyperparameter configurations to model performance, the so-called response surface, behaves as if it only has a small number of relevant dimensions, even when many hyperparameters are nominally present. In practice, only a small subset of hyperparameters substantially impacts the generalization error, while the others leave the response surface nearly flat and can be varied freely without consequence [41]. Because Grid Search forces allocations across all dimensions equally, it wastes the vast majority of its computational budget exploring irrelevant parameters. Random Search, by contrast, tests a unique value for every parameter in every trial, allowing it to thoroughly explore the relevant subspaces with far fewer evaluations [41].

Despite their differences in candidate selection, both Grid Search and Random Search share a fundamental methodological limitation: they are entirely uninformed algorithms [40]. Every trial in their search process is independent, meaning neither method utilizes the performance results of previously evaluated configurations to influence the selection of future points [40]. Consequently, they inevitably waste massive amounts of time and computing resources evaluating poorly performing areas of the configuration space without learning from past failures [40]. This inherent inefficiency establishes the primary motivation for adopting sequential, history-dependent strategies such as Bayesian Optimization.

2.7.2 Bayesian Optimization

To overcome the limitations of uninformed search strategies, Bayesian Optimization (BO) offers a sequential, model-based approach that utilizes the history of previously evaluated configurations to intelligently determine the next hyperparameter values to test [40]. BO operates using two primary components: a probabilistic *surrogate model* that approximates the true, computationally expensive objective function, and an *acquisition function*, such as Expected Improvement (EI), which selects the optimal next point for evaluation [40, 42]. The acquisition function explicitly balances the fundamental optimization trade-off between *exploration* (sampling from under-explored regions of the parameter space where uncertainty

is high) and *exploitation* (sampling from regions already known to yield promising results) [40].

While traditional BO relies on Gaussian Processes (GPs) as surrogate models to directly estimate the conditional probability $p(y|x)$ of the objective value y given the hyperparameters x , the Tree-structured Parzen Estimator (TPE) takes a fundamentally different approach [42]. Instead of asking "given these hyperparameters, what performance can I expect?", TPE inverts the question: "given that a trial turned out well or poorly, what kinds of hyperparameter values were likely to produce it?" [42]. Formally, this corresponds to modeling $p(x|y)$, the distribution of hyperparameter configurations x conditioned on the observed performance y , rather than $p(y|x)$ directly. TPE implements this by splitting all past trials at a performance threshold γ into two groups and fitting a separate density to each: $\ell(x)$ captures the distribution of configurations that performed well, and $g(x)$ captures those that performed poorly [40, 42]. The next candidate is then chosen to maximize the ratio $\ell(x)/g(x)$, favouring configurations that are likely under the good density and unlikely under the poor one [42].

This inverse modeling formulation grants TPE a significant advantage in complex configuration spaces. Because the Parzen estimators are organized in a tree structure, TPE naturally handles conditional hyperparameters, that is, parameters that only exist or only matter depending on the value of another parameter [40]. For example, a dropout rate is only relevant if dropout is enabled at all. If it is disabled, the dropout rate has no effect on performance regardless of its value. A standard GP treats all hyperparameters as independent continuous dimensions and cannot represent such if-then relationships. TPE's tree structure, by contrast, can branch on a parent hyperparameter's value and only then consider its dependent children, faithfully reflecting the actual structure of the search space [40, 42].

Furthermore, TPE offers a substantial computational advantage over both GPs and exhaustive search methods. While GP-based models scale cubically with the number of past observations $\mathcal{O}(n^3)$, making them prohibitively slow as the search history grows, TPE achieves a much lower time complexity of $\mathcal{O}(n \log n)$ [40]. By sequentially directing the search only toward the most promising regions and maintaining low computational overhead per trial, sequential strategies like TPE bypass the exponential combinatorial explosion that plagues Grid Search, enabling the efficient optimization of complex machine learning architectures with many hyperparameters [40].

Table 2.2: Comparison of HPO strategies (n : trials, d : hyperparameters).

Method	Strategy	History	Complexity
Grid Search	Exhaustive	No	$\mathcal{O}(n^d)$
Random Search	Random sampling	No	$\mathcal{O}(n)$
Bayesian (GP)	Sequential, model-based	Yes	$\mathcal{O}(n^3)$
TPE (Optuna)	Sequential, model-based	Yes	$\mathcal{O}(n \log n)$

2.7.3 Optuna as a Practical Framework

Optuna is a modern hyperparameter optimization framework designed around a *define-by-run* API [43], which allows the search space to be constructed dynamically during runtime rather than requiring a static definition beforehand. Unlike traditional tools that require the search space to be statically defined prior to execution, Optuna constructs the hyperparameter search space dynamically within the objective function during runtime, accommodating complex dependencies without relying on rigid external configurations. Optuna utilizes the TPE as its default sampling algorithm, allowing for efficient exploration of these dynamically generated spaces. To ensure reproducibility, the stochasticity of the TPE sampler can be controlled by fixing a global random seed.

Beyond efficient sampling, Optuna accelerates the optimization process through automated trial pruning, which refers to the early termination of unpromising training runs based on intermediate results. This mechanism, formally referred to as *pruning*, periodically monitors the intermediate learning curves of iterative algorithms and preemptively terminates trials that fall below a predefined performance threshold, such as the median performance of previously completed trials. By allocating computational resources exclusively to the most promising configurations, pruning drastically reduces the overall runtime [43]. The specific details regarding whether and how pruning is applied are deferred to Section 5.4.

Ultimately, the adoption of an efficient framework like Optuna is not merely a mathematical convenience, but a fundamental Green AI strategy. The total computational and ecological cost of a machine learning result grows linearly with the number of hyperparameter trials executed [1]. By utilizing Bayesian Optimization to strategically minimize the number of required evaluations, and employing pruning to terminate poor configurations early, the overall energy consumption and resulting carbon footprint of the entire model development process are significantly reduced.

2.8 Chapter Summary

This chapter established the theoretical foundations for the empirical study. The Green AI framework motivated the treatment of ecological cost as a first-class evaluation criterion alongside predictive accuracy. The energy and emissions section formalized how power draw translates into carbon footprint through grid carbon intensity, established why TDP-based estimation is unreliable, and introduced CodeCarbon as the standardized measurement tool. The algorithms section described the five classifiers under examination, their inductive biases on tabular data, and the conditions under which GPU acceleration is genuinely energy-efficient rather than merely faster. Dataset scaling and learning curves established the sub-linear relationship between data volume and performance gains, reinforcing the Green AI argument that ecological costs can outpace accuracy returns. The evaluation metrics section introduced Weighted F1 as the primary classification quality metric and outlined the theoretical limitations of composite efficiency metrics. Cross-validation was formalized as the estimation procedure, with stratification linking directly to the metric choice under

class imbalance. Finally, hyperparameter optimization was covered from exhaustive search through Bayesian Optimization to Optuna, with HPO trial count identified as an ecological cost in its own right. These concepts collectively provide the vocabulary and methodology for the experimental design presented in Chapter 4.

Related Work

This chapter surveys empirical studies that have measured the energy consumption and carbon emissions of machine learning systems. The focus is on what prior work has concretely found, not on the underlying concepts, which are established in Chapter 2. Each section closes by identifying what prior studies leave unmeasured or unresolved, collectively motivating the specific design choices of this study.

3.1 Empirical Energy Benchmarks for Machine Learning

As empirical benchmarking has replaced theoretical complexity analysis as the primary mode of model evaluation, a body of literature has emerged that measures the actual energy and carbon cost of training machine learning systems. This literature has, however, grown unevenly: the studies that sparked the Green AI movement focused almost exclusively on large-scale deep learning, while the empirical benchmarks covering tabular data and classical algorithms ignored energy entirely.

The case for treating energy as a first-class evaluation metric was made most forcefully by early measurements of deep learning workloads. Strubell et al. [3] were among the first to directly measure hardware power interfaces during NLP model training, finding that a single large transformer architecture could emit as much carbon as five cars over their entire lifespans, and helping to establish the Green AI research agenda. Follow-up work examined the structural drivers of this footprint. Dodge et al. [8] expanded the empirical scope to 16 cloud computing regions, demonstrating that hardware choice and grid carbon intensity are as consequential as algorithmic decisions, with the same training run emitting between 7,000 and 26,000 grams of CO₂ depending on the data center. Desislavov et al. [39] extended this perspective to the inference phase, showing that the operational footprint of deployed models can rival the cost of training. A step toward more modest workloads was taken by Henderson et al. [6], whose experiment-impact-tracker framework demonstrated that everyday reinforcement learning experiments on standard hardware generate clearly measurable emissions, establishing ecological efficiency as a concern for ordinary research rather than exclusively for large-scale industrial training. Despite this broadening scope, the entire strand of Green AI benchmarking has remained focused on neural networks trained on unstructured data such as text and images.

The empirical literature on tabular data took a different path entirely. A series of large-scale benchmarks asked which model family generalizes best to structured, heterogeneous datasets and consistently found that classical algorithms hold their own against deep learning. Grinsztajn et al. [21] evaluated 45 tabular datasets and found tree-based models such as Random Forest and XGBoost to outperform deep neural networks, a result echoed

by Gorishniy et al. [15] in a separate benchmark that reaffirmed the dominance of gradient boosting on this data type. Both studies, however, treated energy as entirely outside their scope: runtime was noted as a secondary metric, but physical power consumption and carbon emissions went unmeasured.

Only recently have a small number of studies begun to bridge these two strands. Rodriguez et al. [9] evaluated software-based energy trackers against physical hardware meters, measuring classical machine learning algorithms alongside deep learning models as part of their methodology. Verdecchia et al. [2] took a data-centric approach, empirically measuring how reducing data points and features affects the training energy of six AI algorithms and finding that data-centric dataset modifications can substantially reduce training energy, though the accuracy cost varies depending on whether data points or features are reduced. These contributions are important, but they remain limited to datasets of a few thousand instances and do not examine how energy and accuracy co-evolve across the orders-of-magnitude variation in dataset size that characterizes real-world applications.

The existing literature therefore reveals a clear gap at the intersection of two well-developed but isolated fields. Where Green AI benchmarks have established energy as a measurable, consequential quantity, they have done so exclusively for deep learning on unstructured data [3], [8]. Where empirical benchmarks have rigorously evaluated classical algorithms on tabular data, they have focused entirely on predictive accuracy [15], [21]. A systematic analysis of how energy consumption, carbon footprint, and predictive accuracy jointly scale with dataset size for classical classification algorithms remains absent. The present thesis addresses this gap by empirically investigating these efficiency-performance trade-offs across varying tabular dataset sizes, aiming to identify where simpler, less energy-intensive models become ecologically and practically advantageous.

3.2 Dataset Size Effects on Performance and Energy Consumption

The tension between data volume and model quality has a specific shape. Schwartz et al. [1] conceptualize this in their Equation of Red AI, which models the computational cost of a result as growing linearly with training dataset size. At the same time, accuracy does not grow linearly with data: empirical learning curves consistently follow an inverse power law, meaning that performance gains become increasingly marginal as the training set expands [36]. The practical consequence is that energy costs scale proportionally with data volume while accuracy returns diminish, making large datasets progressively less efficient investments. Koshute et al. [36] formalize the point where these curves cross as the Sufficient Training Set Size (STSS): the minimum volume beyond which additional data is ecologically unjustifiable. A complementary finding from Althnian et al. [37], based on medical classification tasks, adds that the key factor is not absolute size but representativeness, meaning that a small, well-sampled dataset can generalize as well as a much larger but biased one.

The energy implications of dataset reduction have been directly and concretely measured by Verdecchia et al. [2], who trained six AI algorithms while systematically reducing both data points and features in 10% increments. Two distinct patterns emerge from their results.

Reducing the number of features yields energy savings of up to 76% while largely preserving predictive performance, making it a near-cost-free optimization. Reducing the number of data points produces larger savings, up to 92.16% for Random Forest, but at a meaningful accuracy cost that must be weighed explicitly. The effect is also algorithm-specific: KNN achieves a maximum reduction of only 61.72% from fewer rows, while ensemble methods such as Random Forest are considerably more sensitive to training set size. Beyond dataset modifications, the study finds that switching to a less complex algorithm can alone yield energy savings of up to two orders of magnitude, establishing algorithm selection as one of the most powerful levers for ecological efficiency.

These findings rest, however, on a single dataset of 5,574 instances, run exclusively on a standard quad-core CPU and without GPU-accelerated or deep learning baselines. Whether the same energy-accuracy patterns hold at the scale of millions of instances, and whether they generalize across fundamentally different model families, remains empirically unexplored.

3.3 CPU vs. GPU Efficiency: Empirical Evidence

Whether GPU training is more energy-efficient than CPU training cannot be answered by measuring power draw alone. A GPU that draws five times the wattage of a CPU can still consume less total energy, if it finishes the same training task more than five times faster. The relevant quantity is energy, specifically power integrated over time, and how that relationship plays out depends heavily on the algorithm and the dataset.

The power side of this equation is well established empirically. Empirical measurements from cloud training environments report GPU draws of 250 W to 350 W under load, compared to 10 W to 150 W for CPUs [8]. During BERT training, a workload representative of high-utilization deep learning, the GPU alone accounted for 74% of total system electricity. For neural network architectures of this type, the elevated power draw is more than offset by execution speed: the parallelism of deep learning maps efficiently onto GPU hardware, and training that would take days on a CPU completes in hours. For this class of workload, the GPU is the more energy-efficient device.

Classical machine learning algorithms on tabular data present a different picture. Mitchell and Frank [27] measured GPU acceleration for XGBoost specifically, finding speedups of $3\times$ to $6\times$ over a four-core desktop CPU and $1.2\times$ over a dual-socket server with 24 CPU cores combined. These gains are real but considerably smaller than those achieved for neural networks. Given the GPU's substantially higher power draw, a $3\times$ to $6\times$ speedup may not be enough to recover the additional energy cost, and Mitchell and Frank measured only execution time, not energy consumption.

The question of where the break-even point lies for algorithms like XGBoost therefore remains open. A dataset may exist below which the CPU is the ecologically preferable device and above which the GPU's speedup compensates for its power draw, but this has not been empirically established for gradient boosting on tabular data.

3.4 Measurement Tool Accuracy: Software Estimation vs. Hardware Sensors

As Green AI research has grown, so has demand for accessible tools that quantify the energy and carbon cost of model training. The resulting ecosystem of software trackers and analytical estimators has made emissions measurement more tractable, but a growing body of empirical work reveals that different tools applied to the same experiment often report strikingly different results, with the direction of error being systematic rather than random.

The available measurement approaches fall into three categories [9]. Physical external power meters (EPMs) at the wall outlet capture total system consumption but cannot isolate individual processes. On-chip sensors such as Intel’s RAPL and NVIDIA’s NVML offer process-level readings at fine granularity but require specific operating system permissions. Software estimation models infer energy from analytical heuristics such as Thermal Design Power (TDP) or operation counts. In practice, EPMs are expensive to deploy at scale and sensor access is often unavailable, so many researchers rely on estimation-based tools, despite evidence that these introduce substantial and directional errors.

How large these errors can be is documented by Bannour et al. [10], who compared six tracking tools on the same NLP training runs, including ML CO₂ Impact, CarbonTracker, experiment-impact-tracker, and Green Algorithms. Tools relying on analytical models and user-supplied parameters systematically overestimated consumption, while those reading directly from hardware sensors produced more consistent results. The underlying reason for this divergence is that analytical approaches cannot capture the dynamic fluctuations of real hardware. Henderson et al. [6] demonstrated this empirically: TDP-based estimates significantly over- or underestimate actual energy compared to sensor readings, because they ignore real-time utilization shifts and dynamic memory effects, and FLOPs-based estimates often show little correlation with actual consumption across architectures due to hardware-specific firmware behaviors. Rodriguez et al. [9] confirmed this pattern in practice, finding that purely analytical tools consistently deviated from measured ground truth, while sensor-based tools remained closer, though still subject to platform-specific constraints.

More recent work has moved toward direct validation of widely adopted trackers against physical ground truth. Fischer [44] compared CodeCarbon against external power meters across a range of workloads, finding that CodeCarbon consistently underestimates total system power draw by 20 to 30%, as software profilers cannot account for cooling overhead or power supply inefficiencies. Critically, this underestimation is substantially worse for CPU-only workloads than for GPU workloads, with relative errors exceeding GPU-equivalent errors by more than an order of magnitude in some configurations, a finding with direct consequences for any study benchmarking classical machine learning, which runs predominantly on CPU.

3.5 Seasonal and Geographic Carbon Intensity Effects

The carbon footprint of a machine learning experiment is not determined solely by algorithmic and hardware choices. Because electricity grids draw on different fuel mixes depending on geography, time of day, and season, two identical training runs executed on different infrastructure can produce vastly different emissions, making geographic and temporal factors a distinct dimension of Green AI research.

The geographic component of this variability has been measured at scale. Lacoste et al. [7] show that a server in a hydroelectric region such as Quebec emits approximately 20 gCO₂eq/kWh, while a server in a fossil-fuel-dependent region like Iowa emits up to 736.6 gCO₂eq/kWh, a factor of nearly 40 for the same computational workload. Dodge et al. [8] confirmed this pattern across 16 cloud computing regions, finding that choosing the most carbon-efficient data center can reduce training emissions from 26,000 to 7,000 grams of CO₂. Tools such as Green Algorithms [11] and the ML CO₂ Impact Calculator [7] were developed to help practitioners account for these regional differences at the planning stage, making workload placement a practical lever for emissions reduction.

Temporal variation tells a more nuanced story. Dodge et al. [8] find that diurnal variation is meaningful: starting a training run at midnight rather than during peak solar generation hours can increase emissions by up to 8% in certain regions, and they propose time-shifting or dynamically pausing workloads based on real-time grid data. Seasonal variation, however, proves negligible across their 12-month tracking of globally distributed cloud instances, leading them to conclude that month-to-month emissions differences within a single cloud region are small.

This conclusion may not transfer to local electricity grids with a high share of seasonally dependent renewables, which can exhibit considerably stronger swings than global cloud averages suggest.

3.6 Hyperparameter Tuning Efficiency

To achieve state-of-the-art predictive performance, machine learning models require extensive hyperparameter tuning. Since this process involves training architectures multiple times across various configurations, a critical tension emerges between the ecological cost of finding the optimal configuration and the performance benefits it yields.

The Red AI paradigm has normalized evaluating only the final, best-performing model, a practice that obscures the computational budget spent during the development and tuning phases [1]. Empirical measurements reveal that this hidden cost can be substantial. Strubell et al. [3], in an empirical case study tracking the development of a Natural Language Processing pipeline, quantified this overhead directly: a single tuning cycle evaluating 24 configurations multiplied the computational cost of the baseline model by a factor of 24, and the complete development process, which required training 4,789 model variants to find the optimal hyperparameters, consumed 239,942 GPU hours in total. This demonstrates that the tuning phase can multiply the baseline energy footprint of a single model by several orders of magnitude.

To reduce these impacts, researchers must move away from exhaustive methods like grid search. Lacoste et al. [7] explicitly state that grid search is highly inefficient in terms of both model performance and environmental impact, and that replacing it with random or model-based search strategies can accelerate convergence and directly reduce carbon emissions. Empirical benchmarks support this: Bergstra and Bengio [41] show that as few as 8 to 32 random trials can consistently match or outperform exhaustive grid search depending on the size of the search space. Going further, Bergstra et al. [42] demonstrate that parallelizing Bayesian optimization with TPE across five GPUs allowed optimal configurations to be found within 24 hours of wall time, a result that would be unachievable with exhaustive enumeration at comparable computational cost.

While the literature clearly documents the empirical tuning costs for deep learning models in natural language processing and computer vision [3], there is a lack of equivalent benchmarks for classical machine learning on tabular data. How much energy and CO₂ can be saved by reducing the number of tuning trials for algorithms like Random Forest or XGBoost, and at what trial count the marginal accuracy gain ceases to justify the ecological cost, remains unquantified. The present thesis addresses this gap by tracking the carbon emissions of Bayesian hyperparameter optimization via Optuna for Random Forest, XGBoost, and MLP on tabular datasets, providing an empirical basis for evaluating whether the ecological overhead of tuning is proportionate to the performance gains it yields for classical machine learning algorithms.

3.7 Lifecycle Emissions: The Ecological Cost of Deployment and Inference

While the machine learning community has heavily scrutinized the energy and carbon footprint of the training phase, this narrow focus risks obscuring the true ecological cost of an AI system's complete lifecycle. A critical tension thus emerges: does the cumulative energy consumed during the deployment phase eventually outweigh the initial, highly-publicized training footprint?

Industry-scale empirical studies demonstrate that the inference phase is often the dominant factor in a model's total energy consumption. As models are integrated into everyday applications, they are queried millions of times, creating a multiplicative effect where the ongoing energy demand of continuous deployment quickly eclipses the one-time cost of training [39]. This dominance is reflected in industry estimates cited by Patterson et al. [45], who report that both NVIDIA and Amazon Web Services estimated that 80 to 90% of total machine learning computing demand in large-scale deployed systems is attributable to inference. Wu et al. [46] provide further empirical support at the production level, finding that for deployed language models at Meta, the inference phase accounts for 65% of total lifecycle resource consumption, significantly outweighing the training phase.

Consequently, recent research has begun explicitly measuring the inference footprint of pre-trained models at deployment scale. By systematically benchmarking energy consumption across diverse model architectures, Luccioni et al. [47] found that multi-purpose

generative models require orders of magnitude more energy per inference than task-specific systems, even when controlling for the number of model parameters. They also formalize the concept of a cost parity point: the number of predictions at which cumulative inference emissions surpass the initial training footprint, providing a concrete metric for lifecycle analysis. Accurately evaluating this break-even requires accounting for hardware specializations, as Desislavov et al. [39] note that inference efficiency depends heavily on whether a model can fully exploit the underlying accelerator’s optimizations.

Despite the growing acknowledgment of inference costs, the current literature exhibits a significant blind spot. The lifecycle analyses and inference benchmarks provided by these studies rely almost exclusively on complex deep learning architectures applied to unstructured data: Luccioni et al. [47] evaluate only text and image modalities, Desislavov et al. [39] focus on ImageNet and GLUE benchmarks, Patterson et al. [45] restrict their analysis to large NLP models, and Wu et al. [46] examine deep learning recommendation and language models. For classical machine learning algorithms operating on tabular data, direct empirical evidence on lifecycle emissions is near-absent. It therefore remains unquantified at what number of predictions the cumulative inference CO₂ footprint of a model such as Logistic Regression or Random Forest surpasses its training footprint, and how this lifecycle scaling compares across algorithms of varying complexity. The present thesis addresses this gap directly by performing a dedicated lifecycle analysis that systematically scales predictions to empirically measure these break-even points for classical classification algorithms on tabular data.

3.8 Composite Efficiency Metrics: Balancing Performance and Ecological Cost

The singular focus on predictive accuracy in machine learning has historically obscured the underlying computational costs of model training, raising a critical methodological question: how can researchers systematically evaluate both model quality and environmental footprint within a single, standardized framework?

To counteract the tendency to achieve marginal accuracy gains through massive computational expenditure, Schwartz et al. [1] argue that efficiency must be elevated to a primary evaluation criterion alongside accuracy. Measuring algorithmic progress exclusively through predictive performance is fundamentally incomplete. Hernandez and Brown [48] demonstrate that true progress must be quantified by holding performance constant and measuring the reduction in computational cost required to achieve it. Taken together, these theoretical arguments advocate for composite metrics that mathematically integrate compute costs with model capabilities to provide a holistic view of algorithmic efficiency.

In practice, the machine learning community has already recognized the necessity of such combined measurements. Coleman et al. [49] introduced composite metrics such as Time-to-Accuracy and Cost-to-Accuracy in the DAWNbench benchmark, demonstrating that the community accepts end-to-end metrics that evaluate the time and financial budget required to reach a predefined quality threshold. Building upon this precedent, the

MLPerf benchmark suite formally adopted Time-to-Train as its primary evaluation metric, deliberately shifting the focus away from raw accuracy towards the holistic efficiency of the training system [50].

While these benchmarks validate the use of composite metrics, their efficiency components rely exclusively on execution time or financial cost [49], [50]. As established in previous sections, neither time nor cost necessarily correlates with the actual carbon footprint, particularly when hardware power draw and regional grid carbon intensities vary. Recent theoretical frameworks have begun to recognize this shortcoming by proposing conceptual metrics, such as the Job Quality per Cost Rate (JQCR), which aims to balance the quality of a computing job against its sustainability costs [5]. However, this approach remains strictly conceptual and lacks a defined mathematical formula linking specific accuracy measurements to carbon emissions [5]. To date, no established, mathematically operationalized composite metric explicitly incorporates a direct ecological component alongside predictive accuracy. The Efficiency-Weighted F1 (EWF1) metric introduced in this thesis addresses this exact gap: by mathematically integrating a model's F1-score with its operational CO₂ emissions, EWF1 builds upon the accepted paradigms of DAWNBench and MLPerf to provide a composite score explicitly dedicated to measuring ecological efficiency.

Methodology

This chapter outlines the design decisions underlying the empirical evaluation of this study. It describes the research questions, the selected datasets and models, the evaluation metrics, the validation strategy, and the hyperparameter tuning approach that together form the methodological foundation for the experiments conducted in Chapter 5.

4.1 Research Questions

The main research question that this study aims to answer is: *How does the trade-off between energy consumption and predictive performance vary across Logistic Regression, Random Forest, XGBoost and MLP models when the complexity of the models and the size of the dataset are systematically varied?* The following sub-questions will also be addressed in the course of this thesis:

- Under what conditions do the resource consumption of more complex models exceed the achievable accuracy gain over simpler baselines?
- How does energy consumption and CO₂ output scale with increasing neural network complexity (number of layers and neurons)?
- How does energy consumption differ between CPU-based and GPU-based training for equivalent tasks?
- How does reducing the number of training instances on large datasets affect model accuracy and energy consumption across different algorithms?
- At what prediction volume does the cumulative inference CO₂ footprint of a model surpass its training footprint, and how does this break-even point differ across model families?

4.2 Datasets

Dataset size is a primary driver of both computational demand and energy consumption during model training [9]. To evaluate how this relationship varies across different scales, three classification datasets were selected that span several orders of magnitude in size. Classification was chosen as the task type since it is well-established across all selected models and allows for direct comparison of predictive performance via standardized metrics such as Accuracy (ACC) and Weighted F1-Score (WF1). The datasets are summarized in Table 4.1.

Table 4.1: Overview of datasets used in this analysis

Dataset	Instances	Features	Classes	Scale
Wine [51]	178	13	3	Small
Default of Credit Card Clients [52]	30,000	23	2	Medium
HIGGS [53]	11,000,000	28	2	Large

The Wine dataset contains 178 instances across 13 features, where each feature describes a chemical property of wines from the same region in Italy, derived from three different cultivars.

The Default of Credit Card Clients (hereafter: Credit) dataset contains 30,000 instances across 23 features and focuses on predicting payment default among credit card clients in Taiwan.

The HIGGS dataset is the largest dataset used in this study, comprising 11,000,000 instances across 28 features, seven of which are high-level features derived by physicists to help discriminate between the two classes [54]. The classification goal is to distinguish between signal events produced by Higgs bosons and background noise processes.

These three datasets vary significantly in size, enabling a systematic evaluation of how model performance and energy consumption scale across different levels of data complexity.

4.3 Models

Selecting models of varying complexity is central to answering the research questions of this study. A model that achieves competitive accuracy with significantly lower energy consumption represents a more ecologically efficient choice, but this trade-off can only be quantified if the comparison spans a meaningful range of complexity. For this reason, four classification models were selected, ranging from a linear baseline to tree-based ensemble methods and a neural network, covering low, medium, and high computational demand respectively.

Logistic Regression (LR) serves as the lightweight baseline in this study. Despite its simplicity as a linear classifier, it often achieves competitive ACC, making it well suited to quantifying whether the additional resource consumption of more complex models yields a meaningful gain in predictive performance.

Random Forest (RF) was selected as a mid-complexity representative of tree-based ensemble methods. Relevant to this study, Random Forest does not support GPU-accelerated training and is therefore run exclusively on the CPU.

Building on the tree-based approach, Extreme Gradient Boosting (XGB) was chosen as a second mid-complexity model, with the key distinction that it also supports GPU-accelerated training. This makes it possible to directly compare the energy consumption and training time of CPU- and GPU-based tree ensemble methods under identical conditions, which is central to one of the research questions of this study. The suitability of GPU-accelerated XGBoost for datasets at this scale has been demonstrated by Mitchell and Frank, who show that the histogram-based GPU implementation can process a dataset of 10 million instances

with 28 features entirely within GPU memory [27].

The MLP was selected as the most complex model, representing the Neural Network (NN) category. Its highly configurable architecture, particularly the number of hidden layers and neurons per layer, makes it well suited for the architectural complexity experiments conducted in this study, where these parameters are systematically varied to assess their impact on predictive performance and energy consumption.

Table 4.2: Overview of classification models used in this study

Model	Type	Device	Role
LR	Linear	CPU	Baseline
RF	Ensemble	CPU	Mid-complexity
XGB	Gradient Boosting	CPU & GPU	Mid-complexity
MLP	NN	GPU	High-complexity

Together, the selected models cover a broad spectrum of algorithmic complexity, from a linear baseline to gradient boosting and deep learning, enabling a meaningful comparison of both predictive performance and energy efficiency across fundamentally different model types. The following section defines the metrics used to evaluate these dimensions.

4.4 Evaluation Metrics

To answer the research questions of this study, two categories of metrics were employed: classification performance and resource consumption. Accuracy and Weighted F1-Score were selected to measure predictive performance, as they provide a standardized basis for comparing models across datasets of varying class distributions. Energy consumption in kWh and CO₂ emissions in kgCO₂eq were chosen as the primary resource metrics, directly reflecting the ecological cost of model training.

Training time in seconds was included as a secondary resource metric, as it captures computational demand independently of hardware-specific power draw. Furthermore, inference time was measured across all experimental setups to account for the operational footprint. In large-scale deployment scenarios, the cumulative energy demand of repeated predictions can often rival or even exceed the initial training costs [39]. Together, these metrics enable a direct assessment of whether gains in predictive performance justify the associated increase in resource consumption.

Inspired by the efficiency framing proposed in Schwartz et al., a custom composite metric was developed for this study, named the Efficiency-Weighted F1 (EWF1). Equation (4.1) defines this metric, where the numerator is Weighted F1-Score:

$$\text{EWF1} = \frac{\text{WF1}}{1 + \lambda \cdot \hat{c}} \quad (4.1)$$

where $\hat{c}$ denotes the min-max normalized CO₂ emissions per dataset and λ controls the relative importance assigned to ecological cost. A higher λ shifts the metric towards prioritizing carbon efficiency over predictive performance, while a lower λ reduces the

carbon penalty. This allows practitioners to adapt the metric to their specific deployment priorities. All results in this study use the default value of $\lambda = 1$. Higher scores indicate a more favorable trade-off between predictive performance and ecological cost.

It should be noted that $\hat{c}$ is inherently location-dependent, as CO₂ emissions are tied to the carbon intensity of the local electricity grid. Consequently, EWF1 scores reflect the ecological efficiency of the specific experimental setup used in this study and may not generalize directly to other energy infrastructures.

4.5 Validation Strategy

To obtain reliable and generalizable performance estimates, a cross-validation strategy was employed rather than a single train-test split.

Specifically, K-Fold cross-validation was applied with $k = 5$ folds, a value widely recommended in the literature as a practical balance between computational cost and estimation stability [14]. Each model is trained and evaluated five times on non-overlapping subsets of the data, and the final performance score is averaged across all folds. This approach is consistent across all models and datasets, ensuring comparability of results.

To ensure reproducibility, a fixed random state was used for all cross-validation splits, model initialization, and data shuffling. The specific value and its technical implementation are described in Section 5.1.2.

4.6 Hyperparameter Tuning

To ensure a fair comparison across all models, hyperparameter tuning was applied to each model prior to evaluation. Without tuning, default parameters tend to favor certain model families over others, potentially skewing the comparison in terms of both predictive performance and resource consumption [55].

Hyperparameter optimization was performed using Optuna, an optimization framework based on Bayesian optimization. It was configured to run 40 trials per model per dataset, striking a balance between tuning thoroughness and resource consumption, as tuning runs contribute to the overall energy budget of this study.

Tuning was performed separately for each dataset, as the datasets differ substantially in size, feature space, and class distribution. A globally tuned model would not represent optimal performance on any individual dataset and would therefore introduce bias into the comparison.

Implementation and Measurement

This chapter describes how the experimental pipeline was built, how measurements were collected, and which experiments were conducted. It covers the hardware and software environment, the data loading architecture, each model implementation, the hyperparameter tuning setup, and the emissions measurement infrastructure including the custom CPU power monitoring solution. The final section defines the individual experiments that make up this study, directly preceding the results reported in Chapter 6.

5.1 Experimental Setup

5.1.1 Hardware Setup

A critical decision in the experimental design was the selection of the computing platform on which all experiments would be conducted. Two options were considered: the university’s shared computing infrastructure or a local machine. The latter was ultimately chosen, as it provides several key advantages for this study. First, it ensures complete reproducibility, since all experiments are executed on identical hardware configuration, eliminating interfering variables that could arise from shared resource conflicts or system variability. Second, it provides consistent availability and direct control over system configuration, which is particularly important given the extended runtime of the large-scale experiments (particularly on the 11M-instance HIGGS dataset). Third, local execution facilitates fine-grained performance monitoring and measurement of energy consumption, a primary objective of this study.

The hardware platform selected for all experiments is a consumer-grade desktop system with the following specifications: an NVIDIA RTX 4070 graphics processing unit (12 GB GDDR6 memory, 5,888 CUDA cores), an AMD Ryzen 7 5700X processor (8 cores / 16 threads, base clock 3.6 GHz) and 32 GB of DDR4 RAM. The system runs Windows 11 with all code implemented in Python 3.13.0.

Table 5.1: Hardware specification of the experimental system

Component	Specification
CPU	AMD Ryzen 7 5700X (8 cores / 16 threads, 3.6 GHz base)
GPU	NVIDIA RTX 4070 (12 GB GDDR6, 5,888 CUDA cores)
RAM	32 GB DDR4
Operating System	Windows 11

The choice of hardware reflects a deliberate trade-off between experimental scope and practical feasibility. The RTX 4070 represents a mid-to-high-end consumer GPU with

sufficient compute capability to execute GPU-accelerated training for XGBoost and MLP models in reasonable time, while remaining accessible to many practitioners. The Ryzen 7 5700X provides adequate CPU performance for baseline models (Logistic Regression, Random Forest) and CPU-based XGBoost training. Together, this configuration enables direct comparison of CPU- and GPU-accelerated training on the same platform, which is central to answering one of the research questions (*How does energy consumption differ between CPU-based and GPU-based training for equivalent tasks?*).

It is important to note that energy consumption and CO₂ emissions measured in this study are specific to this hardware configuration and the German electricity grid, and may not directly generalize to other systems or geographic regions. However, the relative performance comparisons (e.g., the energy efficiency ranking of different algorithms) are more likely to generalize, as they reflect algorithmic properties rather than hardware specifics alone.

5.1.2 Software Environment and Reproducibility

All experiments were implemented in Python 3.13.0. The main libraries used are described below.

Scikit-learn served as the foundation of the evaluation pipeline. It provided the `LogisticRegression` and `RandomForestClassifier` implementations, as well as `KFold` and `cross_validate` for cross-validation, `StandardScaler` for feature normalization, `make_pipeline` for chaining preprocessing and model steps, and `WF1` and `make_scorer` for evaluation metrics.

The XGBoost library was used for the gradient boosting model, as it provides a highly optimized implementation that supports both CPU and GPU training.

For the MLP, PyTorch [56] was used as the deep learning backend. To integrate the PyTorch model into the scikit-learn evaluation pipeline, the Skorch library [57] was used, which wraps PyTorch models in a scikit-learn-compatible interface.

Hyperparameter tuning for all models except Logistic Regression was performed using Optuna, as described in Section 4.6.

Carbon emissions and energy consumption were tracked using CodeCarbon, which is described in detail in Section 5.5.

To ensure reproducibility, a fixed random seed of 42 was used for all data splits, cross-validation folds, and model initialization. A `requirements.txt` file is provided so that the exact library versions used in this study can be reproduced.

The implementation code was developed with the assistance of Claude Code [58], an AI-powered development tool by Anthropic. All generated code was reviewed, tested, and integrated by the author.

The evaluation relies on `KFold` with five splits and shuffling enabled, executed through `cross_validate()`. Performance is measured using a scoring dictionary that tracks both accuracy and weighted F1 score via `make_scorer(f1_score, average="weighted")`. Final scores are the mean across all five folds.

5.2 Dataset Integration

The three datasets used in this study, Wine, Credit, and HIGGS, were introduced in Section 4.2. This section describes how they are integrated into the software architecture, covering the central configuration file, the data loading pipeline, and dataset-specific constraints.

To ensure a consistent foundation across all model scripts, a central configuration file (`config.py`) was created. It serves as a single source of truth for all dataset-related settings, so that each script accesses the same parameters without redundant definitions.

For each dataset, the configuration stores the file path, the name of the target column, and an optional row limit (`nrows`). The row limit was particularly relevant for the HIGGS dataset, which contains 11 million instances, allowing experiments to be run on smaller subsets where needed. Beyond these common fields, the configuration also captures dataset-specific properties. For example, the Credit dataset includes a header row that must be skipped during loading, and the Wine dataset does not contain column names in the source file, which are therefore defined in the configuration directly. Additionally, a label offset is applied to the Wine target column, as its class labels are encoded as 1, 2, and 3 rather than starting at 0, which is required by the classifiers used in this study.

The global constants `RANDOM_STATE` and `CV_FOLDS` are also defined here, ensuring that all scripts use identical values for random seeding and cross-validation, which is essential for the comparability of results.

To read the datasets from their respective files, a custom `load_data` method was implemented. Depending on the dataset, the appropriate loading mechanism is selected. The Wine and Credit datasets are loaded using pandas' `read_csv` function. For the HIGGS dataset, `read_parquet` is utilized. This was a deliberate decision to save storage space by storing the large HIGGS dataset as a Parquet file rather than a standard CSV.

Following the initial loading phase, the HIGGS data is reduced to a representative subset using a sampling approach controlled by the `nrows` parameter. Once the data size is managed, any columns specified in the configuration file are dropped from the respective dataset. The data is then split into the feature matrix X and the target vector y . If an offset is defined in the config, which is currently only the case for the Credit dataset, it is applied as described in the previous section. Finally, the method returns X and y for further use in the subsequent scripts. This process is visualized in Figure 5.1.

Due to the substantial scale of the HIGGS dataset, several constraints were implemented to ensure the stability of the execution environment. Training a Random Forest model on this specific data led to unmanageable processing times. For this reason, the Random Forest algorithm is explicitly bypassed for the HIGGS dataset and the program terminates using `sys.exit(0)` to prevent system crashes. This allows the process to stay within the boundaries of the available hardware resources.




Figure 5.1: Visualization of the data loading pipeline and preprocessing steps within the `load_data()` method.

5.3 Model Implementations

As outlined in Section 4.3, the following classification algorithms were employed to address the research questions mentioned previously. This section therefore details the specific implementation of each model.

5.3.1 Logistic Regression

The first baseline model is a standard logistic regression. To ensure the model processes all features on an equal scale, the algorithm is embedded in a pipeline starting with a `StandardScaler`. This initial scaling step is crucial for the stability and overall performance of the subsequent classification.

The regression is configured to use the SAG solver by setting `solver='sag'`, chosen for its efficiency on large datasets as described in Section 2.3.1. To ensure proper convergence, the maximum number of iterations is capped at 1000 using the `max_iter` parameter.

This model setup deliberately excludes any hyperparameter optimization. Instead, it acts as a stable baseline to evaluate the raw computational time without the extra overhead of complex search processes.

It is worth noting that the SAG solver is single-threaded and runs on a single CPU core, resulting in roughly 10% utilisation on the 8-core system used in this study. The `cross_validate()` call also uses no `n_jobs` parameter, so the five folds run sequentially. Each model in this study runs with its natural parallelisation behaviour. The implications of this for energy comparability are discussed in Section 5.5.

5.3.2 Random Forest

The second baseline algorithm is the Random Forest classifier. This method belongs to the family of ensemble learning techniques. Unlike logistic regression, this algorithm does not require prior feature scaling. To maximize computational efficiency, the execution is fully parallelized across all available CPU cores by setting the `n_jobs` parameter to `-1`. This configuration allows the processor to run at maximum capacity, which directly influences the subsequent energy consumption tracking and overall runtime measurements.

As established in Section 5.2, Random Forest is explicitly excluded from the HIGGS dataset due to unmanageable training times at 11M instances. Including it would produce extreme runtimes without contributing meaningful insight into the scaling behaviour studied here.

To ensure maximum predictive capabilities for the remaining data, the model configurations were optimized using Optuna. The optimization process specifically targeted maximizing the weighted F1 score. For strict reproducibility, the final parameter combinations for the Wine and Credit datasets are documented in Table 5.2.

Table 5.2: Optimized hyperparameters for the Random Forest classifier obtained via Optuna optimization.

Parameter	Wine	Credit
<code>n_estimators</code>	250	495
<code>max_depth</code>	20	20
<code>min_samples_split</code>	15	17
<code>min_samples_leaf</code>	6	3
<code>max_features</code>	<code>sqrt</code>	<code>None</code>

5.3.3 XGBoost

The next tree-based ensemble model is the XGBoost classifier. This model serves as a direct point of comparison to the Random Forest. A primary reason for including XGBoost in this study is that it allows for a direct comparison of computational efficiency between CPU and GPU processing. To achieve this, the configuration dictionaries are kept identical for both executions. The only difference is the hardware device parameter, which is set to use the GPU. This controlled setup ensures a precise evaluation of the energy and emissions generated by the exact same model architecture on different hardware. This setup directly addresses the third research question: *How does energy consumption differ between CPU-based and GPU-based training for equivalent tasks?*

The CPU versus GPU comparison is derived in the plotting pipeline by reading all HIGGS subset rows from `results.csv` and fitting a log-log model separately for each device using `numpy.polyfit`, with the energy crossover point computed analytically from the intersection of the two fitted lines.

Additionally, XGBoost has a specific technical requirement for handling target variables. While the Random Forest algorithm natively adapts to different class structures, XGBoost

requires explicit instructions for its learning objective. As mentioned in Section 4.2, the Wine dataset consists of three distinct categories. Therefore, the objective parameter is set to `multi:softmax`, paired with the `mlogloss` evaluation metric and a specified number of classes. In contrast, the Credit and HIGGS datasets contain only two categories. Consequently, their objective is configured as `binary:logistic` and evaluated using the standard `logloss` metric.

Similar to the previous baseline algorithm, the parameters tailored to each dataset were determined using the Optuna optimization framework. Furthermore, the algorithm processes the unscaled feature matrix X directly. It also uses the globally defined random state parameter to guarantee a deterministic execution across all hardware setups. The resulting configurations used for the final training phase are detailed in Table 5.3.

Table 5.3: Optimized hyperparameters for the XGBoost classifier across all datasets.

Parameter	Wine	Credit	HIGGS
<code>objective</code>	<code>multi:softmax</code>	<code>binary:logistic</code>	<code>binary:logistic</code>
<code>eval_metric</code>	<code>mlogloss</code>	<code>logloss</code>	<code>logloss</code>
<code>num_class</code>	3	N/A	N/A
<code>n_estimators</code>	273	452	406
<code>max_depth</code>	4	3	8
<code>learning_rate</code>	0.0801	0.0411	0.2263
<code>subsample</code>	0.6558	0.7319	0.8910
<code>colsample_bytree</code>	0.7169	0.6901	0.8222
<code>min_child_weight</code>	4	7	3

5.3.4 Multilayer Perceptron

As discussed in Section 4.3, a Multilayer Perceptron was selected to represent the neural network family. The implementation uses PyTorch as the primary deep-learning framework, combined with the Skorch library. Skorch wraps the PyTorch model in an interface that is fully compatible with the scikit-learn ecosystem. This is necessary because the entire evaluation logic, including cross-validation and pipeline construction, relies heavily on scikit-learn.

A dynamic architecture was chosen to ensure the network structure directly reflects the optimal parameters found during the Optuna tuning process. At runtime, the script reads the best parameter set from a JSON file generated during the tuning phase. It extracts the number of hidden layers via `n_layers` and the neuron count for each individual layer via `layer_i`. These values are then assembled into a list called `layer_sizes`. This list is subsequently passed to the `MLPModule` class. This class constructs the full network dynamically upon instantiation, making the architecture entirely data-driven. The specific configurations resulting from this process for each dataset are detailed in ??.

To minimize the overall training time, the implementation checks for CUDA availability using `torch.cuda.is_available()`. The algorithm runs on the GPU if one is present and automatically falls back to the CPU otherwise. PyTorch requires highly specific numeric types, and neural networks react sensitively to implicit type mismatches. Therefore, the

Table 5.4: Resolved layer-wise architecture of the tuned MLP for each dataset. Each hidden block consists of a Linear layer, Batch Normalization, ReLU, and Dropout. Parameter counts include weights, biases, and Batch Normalization scale and shift parameters.

(a) <i>Wine</i>			(b) <i>Credit</i>			(c) <i>HIGGS</i>		
Lyr	In/Out	Par.	Lyr	In/Out	Par.	Lyr	In/Out	Par.
HL 1	13 → 128	2,048	HL 1	23 → 512	13,312	HL 1	28 → 512	15,872
Out	128 → 3	387	HL 2	512 → 64	32,960	HL 2	512 → 256	131,840
			HL 3	64 → 128	8,576	HL 3	256 → 512	132,608
			HL 4	128 → 256	33,536	Out	512 → 2	1,026
			Out	256 → 2	514			
Total		2,435	Total		88,898	Total		281,346

feature matrix X is explicitly cast to `float32` and the target vector y to `int64` before the training begins. Implementation-level optimizations such as AMP were intentionally omitted to maintain comparability across all tested algorithms and avoid hardware-specific implementation artifacts.

Input features are standardized using a `StandardScaler`, embedded together with the neural network in a `scikit-learn Pipeline`. Since Multilayer Perceptrons lack built-in scale invariance, standardizing the inputs to a mean of zero and unit variance is necessary to keep the gradient flow stable. The `NeuralNetClassifier` provided by `Skorch` receives all architecture parameters through a specific naming convention. It uses the `module_*` prefix to forward keyword arguments directly to the constructor of the `MLPModule`. The Adam optimizer and the `CrossEntropyLoss` criterion are used consistently across all datasets. The learning rate and the batch size are taken directly from the tuning results and are therefore tailored to each dataset. A fixed random seed is applied using `torch.manual_seed` to guarantee exact reproducibility across all iterations.

The maximum number of training epochs is set to 200. This value acts as an upper boundary rather than a strict target. An early stopping callback (`skorch.callbacks.EarlyStopping`) monitors the validation loss after each individual epoch. It halts the training process immediately if no improvement is observed for 10 consecutive epochs. This mechanism prevents unnecessarily long training runs and provides an additional layer of protection against overfitting. Consequently, the upper limit of 200 epochs is rarely reached in practice.

5.4 Hyperparameter Optimization

Building on the design outlined in Section 4.6, this section describes how Optuna was integrated into the evaluation pipeline and which search spaces were defined for each model. The tuning logic is encapsulated in three dedicated scripts (`tune_xgb.py`, `tune_rfc.py`, and `tune_mlp.py`), one per tunable model. All three follow the same structural pattern and differ only in the model that is instantiated and the search space that is queried from the trial object. The target dataset is selected via a command-line argument, and the number of trials is read from the `N_TRIALS` environment variable, with a default of 40. This separation between

configuration and code allows individual tuning runs to be triggered without modifying the source files. The search spaces used across all three tuning scripts are summarized in Table 5.5. Integer and continuous parameters define an inclusive range, while categorical parameters list the discrete choices available to the sampler.

Table 5.5: Hyperparameter search spaces defined in the Optuna tuning scripts.

Model	Parameter	Type	Range / Values
Random Forest	n_estimators	integer	[100, 500]
	max_depth	integer	[3, 20]
	min_samples_split	integer	[2, 20]
	min_samples_leaf	integer	[1, 10]
	max_features	categorical	{sqrt, log2, None}
XGBoost	n_estimators	integer	[100, 500]
	max_depth	integer	[3, 8]
	learning_rate	float (log)	[0.01, 0.3]
	subsample	float	[0.6, 1.0]
	colsample_bytree	float	[0.6, 1.0]
	min_child_weight	integer	[1, 10]
MLP	n_layers	integer	[1, 4]
	layer_i (per layer)	categorical	{64, 128, 256, 512}
	dropout_rate	float	[0.0, 0.5]
	lr	float (log)	[10^{-4} , 10^{-1}]
	batch_size	categorical	{1024, 4096, 8192}

After the dataset has been loaded through the shared `load_data()` function described in Section 5.2, each script constructs a `KFold` splitter using the same global constants `CV_FOLDS` and `RANDOM_STATE` that are also applied during the final evaluation. Reusing the identical splitter ensures that the F1 score returned to Optuna is computed under the same validation conditions as the metrics reported later in Chapter 6, eliminating any discrepancy between the score the optimizer maximizes and the score the study ultimately reports.

The core of every script is an `objective(trial)` function that draws a candidate hyperparameter configuration from the trial object, instantiates the corresponding classifier, and returns the mean weighted F1 score obtained from a five-fold cross-validation via `cross_val_score` with `make_scorer(f1_score, average="weighted")`. This is the only output that Optuna is exposed to. Everything else, such as emissions tracking and timing, is handled outside the objective function.

Each study is created via `optuna.create_study(direction="maximize")` and parameterized with a `TPESampler` seeded with `RANDOM_STATE=42`. Seeding the sampler is what makes the trial sequence reproducible across re-runs. Without it, the suggested hyperparameter combinations would differ on every execution, and the resulting tuning emissions would no longer be comparable. The optimization itself is launched through `study.optimize(objective, n_trials=N_TRIALS)`.

To capture the energy footprint of the tuning process itself, the entire `study.optimize` call is wrapped in a `CodeCarbon EmissionsTracker` with a script-specific `project_name`

(e.g. `tune_mlp_higgs`), so that tuning runs are clearly distinguishable from training runs in the `emissions/` directory. The same tracker is used across all three scripts and is configured identically to the trackers used during training, as described in Section 5.5.

The five-fold cross-validation used within the Optuna objective function relies on the same data partitions as the final model evaluation[cite: 576, 578]. While this approach can lead to a slight optimistic bias in the reported performance metrics due to potential data leakage during the configuration search, it ensures that the computational resources and emissions recorded during the tuning phase are directly representative of the model's interaction with the target validation folds. The implications of this choice regarding the Green AI paradigm are further evaluated in the limitations in Chapter 6.

Once a study completes, the best score, the corresponding hyperparameter dictionary, the duration, the emissions reported by the tracker, and the trial count are written to a shared `best_params.json` file via the `save_best_params()` utility. The function performs an upsert keyed by model and dataset, so that re-running a single tuning configuration overwrites only its own entry and leaves the other results untouched. The model scripts introduced in Section 5.3 consume this file through the corresponding `load_best_params()` helper, which decouples tuning from training entirely: the final training pipeline never instantiates Optuna, it merely reads the previously persisted hyperparameter dictionary and passes it to the classifier constructor.

5.5 Emissions Measurement

Measuring energy consumption and CO₂ emissions is the central objective of this study, and the measurement infrastructure therefore required particularly careful design. This section describes how CodeCarbon was integrated into each script, how measurement coverage was defined across model and dataset combinations, and what hardware-specific constraints and limitations apply to the collected values.

Energy consumption and CO₂ emissions were tracked using CodeCarbon as described in Section 2.2.4. In each model script, an `EmissionsTracker` instance is created before the cross-validation loop begins and configured with two key parameters: `output_dir`, which points to the shared `emissions/` directory at the project root, and `project_name`, which uniquely identifies the run. The tracker is then started by calling `tracker.start()` immediately before the `cross_validate()` call, and stopped via `tracker.stop()` once all folds have completed. The return value of `tracker.stop()` is a single float representing the total CO₂-equivalent emissions in kilograms accumulated over the entire tracked interval. This value is what gets passed to `save_results()` and written to `results.csv`.

The placement of the tracker boundaries was a deliberate choice. Data loading and preprocessing happen before `tracker.start()`, so they are excluded from the measurement. Only the cross-validation itself, which includes all five training and evaluation passes, falls inside the tracked interval. This ensures that the reported emissions value reflects the cost of model training and evaluation exclusively, and that values are directly comparable across models and datasets without interference. While standard deployment practices typically conclude with a final retraining phase on the entire dataset, this study deliberately evaluates

the cumulative resource footprint of the cross-validation loop. This approach captures the realistic, unavoidable computational burden of model evaluation and avoids the redundant ecological cost of an additional training pass, as the relative efficiency rankings between algorithms remain structurally unaffected.

Figure 5.2 illustrates this structure. Prior to training, data is loaded and, where required, features are normalized using `StandardScaler`. Scaling was applied to Logistic Regression and MLP, as these models are sensitive to the magnitude of input features, while tree-based models (Random Forest, XGBoost) were trained on the raw input data.




Figure 5.2: Training and emissions tracking procedure per model-dataset combination. Dashed gray boxes indicate CodeCarbon `EmissionsTracker` calls, and dashed blue boxes indicate `CPUPowerMonitor` calls (background thread, 0.25s sampling). Both run in parallel during cross-validation. After training, `compute_corrected_co2` replaces CodeCarbon’s TDP-based CPU estimate with the hardware sensor value.

To allow fine-grained analysis, each model–dataset combination is tracked by a dedicated `EmissionsTracker` instance, identified by a unique `project_name` following the pattern `model_dataset`, for example `lr_wine`, `rf_credit`, or `xgb_gpu_higgs`. This scheme makes individual runs immediately identifiable in the `emissions/` directory, where CodeCarbon writes one CSV file per tracked run in addition to returning the scalar value in code. Tuning scripts follow the same pattern preceded by `tune_`, for example `tune_mlp_higgs` or `tune_xgb_cpu_wine`. In total, the experiment produces tracked emissions records for 14 training runs and up to eight tuning runs. The separate CSV files serve as a raw audit trail, while the values extracted via `tracker.stop()` are the primary source for the analysis presented in Chapter 6.

The general measurement approach used by CodeCarbon, including direct GPU readout via NVML, RAM estimation per DIMM slot, and TDP-based CPU fallback on Windows, is described in Section 2.2.4. The following paragraphs focus on the concrete implications for

the hardware used in this study.

CodeCarbon’s TDP-based CPU estimation on Windows has been replaced in this study by a custom `CPUPowerMonitor` class implemented in `models/power_monitor.py`. It runs in a background thread and polls the CPU Package Power sensor every 0.25 seconds via the `LibreHardwareMonitor` library (`HardwareMonitor` Python package [59]). This library reads the CPU Package Power sensor directly from the hardware, providing true power values rather than a TDP approximation. In practice, CodeCarbon underestimated CPU power by a factor of roughly $9\times$ compared to the hardware sensor (e.g. 6.6 W vs. 57.8 W for Logistic Regression on the Credit dataset). The corrected CO₂ value in `results.csv` therefore replaces CodeCarbon’s CPU contribution with the hardware-measured value, while GPU and RAM energy are still taken from CodeCarbon. The original CodeCarbon total is retained as `co2eq_codecarbon_kg` for reference.

This approach does, however, come with notable limitations. Because `LibreHardwareMonitor` requires privileged access to hardware sensors, all measurements in this study were conducted with the process running under administrator privileges. Without elevated rights, the `CPUPowerMonitor` fails to retrieve sensor data and the corrected CO₂ estimate cannot be computed. Beyond the access requirement, the approach is inherently platform-specific. `LibreHardwareMonitor` is a Windows-only solution, adopted here precisely because Intel RAPL, the interface CodeCarbon uses for direct CPU energy readings on Linux, is not accessible on Windows. The hardware-corrected emissions values reported in this study are therefore tied to this Windows-specific stack and cannot be reproduced directly on other operating systems.

A secondary limitation is that CodeCarbon measures at the process level, meaning background system activity on Windows 11 contributes to the recorded energy. This effect is expected to be small relative to the dominant training workload but cannot be fully eliminated.

A final limitation concerns the differences in CPU parallelisation across models. Logistic Regression, for instance, relies on the SAG solver, which executes on a single core, whereas Random Forest and XGBoost exploit multi-threading across all available threads. This disparity directly affects the instantaneous power draw: a single-threaded workload draws considerably fewer watts. However, whether the reduced wattage is offset by a proportionally longer runtime introduces a confounding variable when evaluating pure algorithmic efficiency. The implications of this trade-off for the comparability of emissions estimates are discussed in Chapter 6.

5.6 Inference Latency Measurement

Inference latency was measured separately after training, outside the emissions tracker. For each model, a single row from the dataset was passed through the trained model repeatedly, and the elapsed time was captured using `time.perf_counter()`. This provides a hardware-level estimate of per-sample prediction latency, which is relevant in deployment scenarios where repeated inference can accumulate significant energy costs [39].

Only a single row was passed to the model rather than a full batch. This was intentional:

batching would measure throughput rather than the actual time it takes to process one request, which is the more relevant figure in deployment. A single-prediction latency can be extrapolated to real-world usage, where a model may handle millions of requests and the per-sample cost accumulates into a meaningful energy expenditure.

The model used for this measurement is taken from the first cross-validation fold via `cv_results["estimator"][0]`, which is made available by passing `return_estimator=True` to `cross_validate()`. Because inference latency is strictly a function of the model's fixed architecture rather than the size of its training set, utilizing the estimator from the first fold provides an exact timing measurement without the computational redundancy of a full retraining phase. To obtain a stable latency estimate, the measurement is repeated 100 times on the same single-row slice `X[:1]`, preceded by one discarded warmup call to eliminate cold-start effects such as JIT compilation or cache misses. The median of the 100 measurements is recorded as the inference time, making it robust to transient system load spikes that would distort a single-shot measurement.

The recorded latency values are subsequently used in a dedicated lifecycle analysis (`run_lifecycle_analysis.py`), which estimates the number of predictions at which the cumulative inference CO₂ footprint of each model surpasses its one-time training footprint. This break-even point is derived by combining the per-sample inference latency with the average CPU power draw recorded during training (`cpu_power_hw_w`) and the carbon intensity of the local grid, which is itself derived from the training measurements. It should be noted that this approach uses training-phase power as a proxy for inference-phase power; actual inference power draw may differ, and this constitutes a methodological limitation acknowledged in Chapter 6.

5.7 Results Storage and Orchestration

The results of every model run are written to `results/results.csv`, which serves as the central data source for all subsequent analysis and visualisation. Each row represents one complete model-dataset combination and is appended to the file rather than overwriting it, so that results from multiple runs accumulate without data loss. The file is managed using Python's `csv.DictWriter`, which writes a header row on first creation and appends subsequent rows without repeating the header. A consequence of the append-only design is that re-running a configuration produces duplicate rows, which must be filtered during analysis. Table 5.6 lists all columns.

The primary emissions value is `co2eq_kg`, which combines the hardware-sensor CPU reading with CodeCarbon's GPU and RAM estimates as described in Section 5.5. The column `co2eq_codecarbon_kg` retains the original CodeCarbon total for reference. The difference between the two columns is examined in Chapter 6.

Inference latency is stored separately in `results/inference_time.csv`, as it is measured outside the emissions tracker and has a different structure. Like `results.csv`, it uses append-only writes via `csv.DictWriter`. Table 5.7 shows its columns.

Four core orchestration scripts coordinate the data collection phase of the experiment. Each model script is launched sequentially as an isolated subprocess rather than being

Table 5.6: Column schema of `results.csv` including component-level energy tracking.

Column	Format	Description
<code>timestamp</code>	<code>datetime</code>	Date and time of the run
<code>model</code>	<code>string</code>	Model name
<code>dataset</code>	<code>string</code>	Dataset name
<code>nrows</code>	<code>int / all</code>	Subset size; <code>all</code> for full dataset
<code>accuracy</code>	<code>float</code>	Mean CV accuracy
<code>f1</code>	<code>float</code>	Mean weighted F1 score
<code>co2eq_kg</code>	<code>float</code>	Corrected CO ₂ (Hardware CPU + CC GPU/RAM)
<code>co2eq_codecarbon_kg</code>	<code>float</code>	Original CC estimate (static 381 g/kWh factor)
<code>cpu_power_hw_w</code>	<code>float</code>	Average CPU package power (W)
<code>cpu_energy_hw_wh</code>	<code>float</code>	Total CPU energy (Wh)
<code>gpu_energy_wh</code>	<code>float</code>	Total GPU energy from CodeCarbon (Wh)
<code>ram_energy_wh</code>	<code>float</code>	Total RAM energy from CodeCarbon (Wh)
<code>training_time_s</code>	<code>float</code>	Total CV wall-clock time (s)

Table 5.7: Column schema of `inference_time.csv`.

Column	Format	Description
<code>timestamp</code>	<code>datetime</code>	Date and time of the run
<code>model</code>	<code>string</code>	Model name
<code>dataset</code>	<code>string</code>	Dataset name
<code>nrows</code>	<code>int / all</code>	Subset size; <code>all</code> for full dataset
<code>inference_time</code>	<code>float</code>	Single-row prediction latency (s)
<code>co2eq_kg</code>	<code>float</code>	Corrected CO ₂ of the training run

imported directly. This is important because CodeCarbon operates at the Python process level: running multiple trackers within the same process risks state leakage between measurements. Subprocess isolation ensures that each tracker starts and stops cleanly in a fresh runtime environment. However, as noted in Section 5.5, the hardware sensors still measure whole-system power draw during this isolated execution, meaning Windows background activity remains part of the final recorded footprint.

`run_all.py` is the main entry point for the full cross-dataset benchmark. It runs in two sequential phases. Phase 1 executes all tuning scripts for Random Forest, XGBoost, and MLP across their respective datasets, saving the best parameters to `best_params.json`. Phase 2 then runs all five model scripts across all three datasets using those saved parameters (except RF on the HIGGS dataset as mentioned earlier).

The remaining three scripts execute specific experimental variations. `run_xgb_breakeven.py` runs XGBoost CPU and GPU directly on nine HIGGS subset sizes to provide data for the hardware efficiency crossover analysis. `run_scaling_all_models.py` runs the remaining algorithms across designated HIGGS subset sizes to capture scaling behavior. Both subset scripts control the dataset size externally via the `TEST_NROWS` environment variable, which the data loader reads to draw a stratified sample. Finally, `run_mlp_variation.py` isolates the effect of neural network complexity by training a grid of 12 MLP architectures on HIGGS, overriding the tuned hyperparameters with neutral defaults to ensure a fair comparison across different depths and widths.

5.8 Experiments

Having established the implementation and measurement infrastructure, this section describes the individual experiments conducted in this study. Each experiment is motivated by one or more of the research questions stated in Section 4.1 and addresses a specific dimension of the energy-performance trade-off.

5.8.1 Main Cross-Dataset Benchmark

The main cross-dataset benchmark forms the empirical backbone of this study. It trains all five models on all three datasets, with the single exception of Random Forest on HIGGS, which is excluded due to prohibitive runtimes at 11 million instances. The resulting 14 model-dataset combinations are evaluated under identical conditions: dataset-specific hyperparameters from the Optuna tuning runs, five-fold cross-validation with a fixed random state, and the corrected emissions measurement pipeline described in Section 5.5. Logistic Regression is the only model not subject to hyperparameter tuning, as its role is that of a lightweight linear baseline whose value lies precisely in its simplicity.

For each of the 14 runs, accuracy and weighted F1-score are averaged across the five folds, while CO₂ emissions and training time are recorded as totals over the full cross-validation loop. Inference latency is measured separately on a single held-out sample after training. From these values, the Efficiency-Weighted F1 score (EWF1) is computed during analysis and serves as the primary composite metric for comparing the performance-to-cost ratio across models and datasets.

The benchmark is designed to support comparison along two axes simultaneously. On the model complexity axis, ranging from Logistic Regression through the tree ensembles to MLP, it exposes at what point additional complexity stops producing accuracy gains that justify the associated energy cost. On the dataset scale axis, spanning three orders of magnitude from Wine to HIGGS, it reveals how the relative standing of models shifts as training data grows. The inclusion of both XGBoost CPU and XGBoost GPU as separate entries also directly feeds into the CPU versus GPU efficiency comparison developed in the scaling experiment.

5.8.2 Dataset Scaling Experiment

The main benchmark trains each model on fixed datasets, which makes it impossible to separate the effect of dataset scale from other algorithmic differences. The scaling experiment addresses this by holding everything constant except the number of training instances. HIGGS is the only dataset large enough to make this meaningful: with 11 million instances, it can be subsampled across a wide enough range for the scaling behaviour to become visible. Wine and Credit are too small to show any trend across multiple subset sizes.

The experiment trains four models (Logistic Regression, XGBoost CPU, XGBoost GPU, and MLP) on seven HIGGS subset sizes: 1,000, 10,000, 100,000, 500,000, 1,000,000, 5,000,000, and 11,000,000 instances. The range starts at 1,000 rows to capture the low-data regime where

models may not yet converge and where resource costs are dominated by fixed overhead rather than computation. The upper endpoint is the full dataset, giving 28 model–subset combinations in total. Consistent with the cross-validation strategy established in Section 2.6, subsets are drawn via stratified sampling to preserve the original class distribution across all sizes. Random Forest is excluded for the same reason as in the main benchmark.

The hyperparameters used at every subset size are those obtained from tuning on the full HIGGS dataset and are not re-optimised per subset. Re-tuning at each size would conflate scale effects with configuration effects and make it impossible to attribute observed differences to dataset size alone. The practical consequence is that at smaller subsets the hyperparameters are over-specified for the data, so performance figures at 1,000 or 10,000 instances reflect a model trained with a configuration designed for 11 million rows, not the best achievable result at that size.

For each combination, CO₂ emissions, training time, and weighted F1-score are recorded using the same measurement pipeline as the main benchmark. Because XGBoost is the only model present in both a CPU and a GPU variant, the scaling data also directly supports the CPU versus GPU comparison: it shows at which dataset size the GPU’s higher power draw starts to be offset by its faster training time, making the energy crossover visible as part of the broader scaling picture rather than as an isolated finding.

5.8.3 MLP Architecture Variation

The architectural variation experiment isolates the effect of MLP complexity on training emissions and predictive performance. The main benchmark uses one tuned architecture per dataset and therefore cannot answer how those metrics scale as layers and neurons are added. This experiment fills that gap by training the MLP on HIGGS across a systematic grid of 12 architectures: three widths (128, 256, and 512 neurons per layer) crossed with four depths (one, two, three, and four hidden layers). Each layer in a given configuration uses the same width, which keeps the design space tractable and ensures the depth and width axes can be interpreted independently of one another. The full grid is shown in Table 5.8 together with the resulting parameter counts on HIGGS (28 input features, 2 output classes), which range across more than two orders of magnitude.

Table 5.8: MLP architectures evaluated in the variation experiment, with trainable parameter counts for HIGGS. Counts include linear layer weights and biases.

Depth	Width 128	Width 256	Width 512
1 layer	3,970	7,938	15,874
2 layers	20,482	73,730	278,530
3 layers	36,994	139,522	541,186
4 layers	53,506	205,314	803,842

HIGGS is the only one of the three datasets where architectural variation has a meaningful effect to measure. Wine has too few instances to support a deep network at all, and Credit is structurally simple enough that even Logistic Regression achieves accuracy comparable to a tuned MLP, making any architecture a wash in performance terms. With

11 million instances and a non-trivial signal-versus-background classification task, HIGGS gives the network enough room to actually distinguish between configurations.

The non-architectural hyperparameters (learning rate, dropout rate, and batch size) are deliberately not taken from the HIGGS tuning results. The Optuna run optimised those values jointly with one specific architecture, so reusing them would silently favour that architecture against the rest of the grid. The experiment instead fixes neutral, architecture-independent defaults: a learning rate of 10^{-3} (the Adam default from Kingma and Ba), a dropout rate of 0.3 (a standard mid-range regularisation value), and a batch size of 4,096 (the middle value of the Optuna search space defined in Table 5.5). This choice may produce slightly suboptimal absolute results for any single configuration, but the relative comparison across the grid becomes considerably more defensible because no architecture starts with a hidden advantage from the tuning phase.

Each of the 12 runs follows the same measurement pipeline as the main benchmark: five-fold cross-validation with the shared random state, CO₂ and training time tracked over the full cross-validation loop, weighted F1 averaged across folds, and inference latency measured separately. The resulting data supports two orthogonal comparisons. Holding width constant while increasing depth reveals how layer stacking contributes to emissions and accuracy. Holding depth constant while increasing width reveals the same for neuron count. Together, these two views directly answer the research question of how energy consumption and CO₂ output scale with neural network complexity.

5.8.4 Temporal Carbon Intensity Analysis

CodeCarbon converts measured energy into CO₂ using a static carbon intensity factor of 381 gCO₂eq/kWh for Germany. This single number averages out two sources of variation that the literature has identified as significant for the German grid: a strong diurnal cycle driven by solar production, and a seasonal swing driven by the winter coal share versus summer renewable share [8]. The static factor is therefore convenient but unrepresentative of the actual emissions of any training run that was not coincidentally executed at the annual mean intensity. This analysis quantifies how much the reported emissions of this study would have differed under realistic timing.

The experiment is purely retrospective and requires no retraining. A full year of hourly carbon intensity data for the German bidding zone is fetched from the ElectricityMaps API [60], covering 8,760 observations across all four seasons and the full diurnal cycle. The actual energy in kWh consumed by every training run is derived from the recorded CO₂ value by inverting the static factor:

$$E_{\text{kWh}} = \frac{\text{CO}_{2,\text{recorded}}}{0.381 \text{ kg/kWh}} \quad (5.1)$$

Each run's energy is then re-multiplied by carbon intensity values drawn from the API data to produce counterfactual emissions under the eight scenarios defined in Table 5.9. These scenarios isolate three different sources of variation: the global average across the full year, the extremes of best- and worst-case timing, and the structurally informed seasonal-

diurnal combinations where renewables or fossil sources dominate.

Table 5.9: Carbon intensity scenarios used in the counterfactual analysis. Each scenario is computed from one year of hourly ElectricityMaps data for the German bidding zone.

Scenario	Definition
Static (CodeCarbon)	Fixed annual factor used by CodeCarbon, 381 g/kWh
Year mean	Mean over all 8,760 hourly observations
Year best hour	Single lowest-intensity hour of the year
Year worst hour	Single highest-intensity hour of the year
Winter mean	Mean over December, January, February
Summer mean	Mean over June, July, August
Summer midday	Mean of summer hours 11–14 UTC (peak solar)
Winter evening	Mean of winter hours 17–20 UTC (peak coal and gas)

Two questions guide the analysis. First, how large is the gap between the static-factor estimate and the realistic best- and worst-case timing for the same workload. If the gap is small, CodeCarbon’s static factor is defensible as a practical approximation. If the gap is large, the reported emissions of this study should be read as one point on a wide distribution rather than as a fixed value. Second, the analysis identifies whether the diurnal or the seasonal axis dominates the variation, which has direct implications for practitioners deciding when to schedule large training jobs on local infrastructure. The script outputs a heatmap of mean intensity by hour-of-day and month, plus a per-run counterfactual CSV with the recomputed emissions for every scenario.

5.8.5 Lifecycle Break-even Analysis

The main benchmark measures the one-time training footprint of each model, but training CO₂ is only part of a model’s total ecological cost. In production, a deployed model is queried repeatedly, and the cumulative inference footprint grows linearly with the number of predictions made. The lifecycle analysis quantifies at what prediction volume this cumulative inference CO₂ surpasses the initial training footprint, and how this break-even point differs across model families.

The analysis is purely post-hoc and requires no additional training runs. It combines two data sources already collected during the main benchmark: the corrected training CO₂ from `results.csv` and the median single-row inference latency from `inference_time.csv`. The CO₂ emitted per single prediction is estimated as:

$$\text{CO}_{2,\text{inference}} = \frac{t_{\text{inference}} \times P_{\text{CPU}}}{3,600,000} \times I_{\text{carbon}} \quad (5.2)$$

where $t_{\text{inference}}$ is the median inference latency in seconds, P_{CPU} is the average CPU package power in watts recorded during training, and I_{carbon} is the carbon intensity in kgCO₂/kWh derived from the training measurements as $I_{\text{carbon}} = \text{CO}_{2,\text{train}}/E_{\text{train}}$. The break-even prediction count follows directly as $N_{\text{break-even}} = \text{CO}_{2,\text{train}}/\text{CO}_{2,\text{inference}}$.

This approach has one deliberate simplification: P_{CPU} is the power draw recorded during training rather than during inference. Inference on CPU-bound models typically draws less

power than training because there is no gradient computation or backward pass. Using training-phase power therefore produces a conservative break-even estimate; the true break-even will occur at a higher prediction count if inference power is lower. This limitation is acknowledged in Chapter 6.

Results and Discussion

6.1 Emissions Measurement Validation

Before presenting the benchmark results it is worth establishing the reliability of the emission values. As described in Section 5.5, CodeCarbon's TDP-based CPU estimate was replaced by direct hardware-sensor readings. Figure 6.1 shows the ratio between the two for every model-dataset combination.




Figure 6.1: CodeCarbon underestimation factor (hardware-corrected / CodeCarbon) per model and dataset. All values exceed 1, confirming systematic undercounting. Higher factors for Wine and Credit reflect CodeCarbon's TDP-based CPU estimation being least accurate on short training runs; RF has no HIGGS measurement.

Zahlen aus Figure kurz benennen: z.B. Median-Faktor, Min/Max

6.2 Hyperparameter Tuning Overhead

Tuning itself consumes energy that forms part of the total experimental budget. Table 6.1 shows the CO₂ cost and duration of every Optuna study (40 trials each), together with the best weighted F1 achieved.

Werte nach finalem Run aktualisieren

Table 6.1: Hyperparameter tuning overhead per model and dataset. Each run comprised 40 Optuna trials optimising weighted F_1 . CO_2 was measured with the hardware-corrected tracker; duration covers the full `study.optimize` call. Random Forest was not tuned on HIGGS; literature defaults were used instead.

Dataset	Model	CO_2 (g)	Duration	Best F_1
Wine	Random Forest	0.48	77 s	0.9831
	XGBoost	1.04	2.1 min	0.9888
	MLP	0.39	46 s	0.9943
Credit	Random Forest	10.51	13.6 min	0.8001
	XGBoost	1.08	2.2 min	0.8008
	MLP	6.34	15.7 min	0.8044
HIGGS	XGBoost	21.9	50 min	0.7597
	MLP	2149	95.4 h	0.7873

6.3 Main Cross-Dataset Benchmark

Figure 6.2 shows the complete benchmark results across all 14 model-dataset combinations. The six panels report CO_2 emissions (corrected), training time, inference latency, accuracy, weighted F_1 -score, and EWF1.

Kernergebnisse in Prosa: konkrete Zahlen für Accuracy, F_1 , Emissions je Dataset nennen. Keine Interpretation, nur Fakten: "LR erreicht X auf Wine, RF Y auf Credit, MLP Z auf HIGGS."

maybe
split
dash-
board
figure
so page
rotation
is not
needed




Figure 6.2: Model comparison dashboard: CO₂ emissions (corrected), training time, inference time, accuracy, F1-score, and carbon-penalised F1 (EWFI) across all models and datasets.

6.4 XGBoost: CPU vs. GPU

Figure 6.3 compares XGBoost CPU and GPU across all three datasets on CO₂ emissions and training time.




Figure 6.3: XGBoost CPU vs. GPU across datasets. CO₂ and training time on log scale.

Konkrete Zahlen: CPU vs. GPU Emissions und Zeit je Dataset. Wo ist CPU günstiger, wo GPU?

6.5 HIGGS Dataset Scaling

Abbildung: Scaling-Plot (CO₂ und F1 über Subset-Größen für LR, XGB-CPU, XGB-GPU, MLP)

Zahlen: Wie skalieren Emissions und F1 mit wachsender Datenmenge? Wo flacht der Accuracy-Gewinn ab?




Figure 6.4: XGBoost break-even point comparison between different subset sizes of the HIGGS dataset.




Figure 6.5: All models evaluated on different scales of the HIGGS subsets.

6.6 MLP Architecture Variation

Abbildung: Heatmap oder Grouped-Bar (Depth x Width vs. CO2 / F1 auf HIGGS)

Zahlen: Wie ändern sich Emissions und F1 mit Tiefe/Breite? Wann hört der Gain auf?

6.7 Carbon Intensity Analysis

The analyses conducted in the preceding sections treat CO₂ emissions as a fixed function of the training run: given the same model, dataset, and hardware, the reported figure does not change. This is a consequence of how CodeCarbon works: it converts measured energy into CO₂ using a static intensity factor of 381 gCO₂eq/kWh for Germany, a calculated average that cannot reflect when a run actually takes place. The German grid, however, draws on a fundamentally different fuel mix in January than in July, and during peak evening demand than during solar-rich midday hours, a structure visible directly in the hourly ElectricityMaps data used in this analysis. If this temporal variation is large enough, then the timing of a training run is itself a meaningful ecological decision, one that requires no changes to code, hardware, or model architecture. This section evaluates that hypothesis directly by counterfactually re-assigning the energy of every training run in this study to different timing scenarios derived from one year of hourly ElectricityMaps data, as described in Section 5.8.4.

The first result is somewhat reassuring for the existing measurements: as shown in Figure 6.6, the static CodeCarbon factor of 381 gCO₂/kWh is a defensible approximation of the annual mean, which the empirical data places at 341 gCO₂/kWh. A workload executed at a uniformly random point during the year would therefore be assigned only a modest systematic error by CodeCarbon. The more important finding is the spread around that average. The best recorded hour of 2024 reaches 115 gCO₂/kWh, and the worst reaches 658 gCO₂/kWh, a factor of 5.7× between two hours. The absolute extremes are not practically exploitable, since no practitioner can guarantee their run coincides with the single cleanest or dirtiest hour of the year. The structurally informed seasonal-diurnal combinations are the actionable range: summer midday hours average 193 gCO₂/kWh while winter evening hours average 389 gCO₂/kWh, a factor of two achievable through deliberate scheduling alone.

Because training energy is a fixed physical quantity for any given run, every percentage difference in grid intensity translates directly into the same percentage difference in CO₂. Figure 6.7 expresses this proportionality relative to the static CodeCarbon baseline, making the stakes concrete. Scheduling during summer mean conditions would reduce the reported emissions of this study by approximately 20 % compared to the static estimate, requiring no modification to the experimental setup whatsoever. The best-hour scenario pushes this reduction to 70 %, while the worst-hour scenario reverses it into a 73 % increase. These percentages are identical for every model and dataset, because the intensity factor scales uniformly; what differs is the absolute magnitude of the effect. For lightweight runs such as Logistic Regression on Wine, the difference between any two scenarios is negligible in absolute terms. For the MLP trained on HIGGS, the gap between best-hour and worst-hour timing amounts to approximately 208 gCO₂, a quantity that exceeds the total training footprint of every Wine and Credit run in this study combined. This asymmetry reveals an important practical principle: grid-aware scheduling is most valuable precisely where algorithmic and hardware efficiency decisions also matter most, namely on large, compute-intensive workloads where the absolute energy expenditure is high enough for timing to




Figure 6.6: Carbon intensity (gCO₂/kWh) for eight timing scenarios

make a meaningful difference.

To understand why the diurnal dimension matters as much as the seasonal one, consider the solar duck curve visible in Figure 6.8 for 24 July 2024. Carbon intensity declines from roughly 300 gCO₂/kWh in the morning as solar generation ramps up, reaches a minimum of 163 gCO₂/kWh near midday, then recovers sharply to a peak of 461 gCO₂/kWh once solar output falls in the evening. The static CodeCarbon factor of 381 gCO₂/kWh intersects this curve only during the early morning and late evening hours, the two dirtiest windows of the day on a summer date with strong solar generation. During the core daylight hours, the static factor overestimates the actual intensity by up to a factor of two. This is not an anomaly specific to this particular date: it is the structural consequence of Germany’s growing solar capacity, which consistently depresses midday intensity throughout the summer months. A practitioner who runs training jobs during office hours on summer days is inadvertently benefiting from this effect, while their CodeCarbon reports attribute a higher carbon cost than actually occurred.

The seasonal dimension adds a second layer of variation on top of the diurnal one. The ten-day windows shown in Figure 6.9 reveal that winter and autumn both sit above the static CodeCarbon factor on average, meaning CodeCarbon systematically *underestimates* winter emissions rather than overestimating them. The winter window averages 474 gCO₂/kWh, nearly 25 % above the static factor, while spring (347 gCO₂/kWh) and summer (361 gCO₂/kWh) fall meaningfully below it. This directional asymmetry matters for transparency: the runs in this study were conducted during periods of varying grid intensity, and the static factor provides no indication of whether a given run’s reported footprint is an over- or underestimate. An equally important finding is that within-season variance is large enough to undermine coarse seasonal scheduling as a reliable strategy. The winter window alone spans from 311 to 614 gCO₂/kWh across ten consecutive days, a range wider than the gap between winter and summer means. A practitioner who decides to




Figure 6.7: Counterfactual CO₂ relative to the static CodeCarbon baseline (= 100 %)




Figure 6.8: Hourly carbon intensity on 24 July 2024 in Germany

schedule training in July rather than January improves their expected footprint, but cannot assume low intensity on any given July day without consulting real-time grid data.




Figure 6.9: Hourly carbon intensity for representative ten-day windows per season, Germany 2025

The year-long heatmap in Figure 6.10 brings both dimensions together and makes their joint structure immediately legible. The lowest-intensity region of the entire year is summer midday: a consistent band of green running from May through August between approximately 10:00 and 15:00, reflecting the hours when solar generation is both abundant and at its seasonal peak. The highest-intensity region is winter evenings: the band of red and orange concentrated between 17:00 and 21:00 from November through February, when solar output has already dropped to near zero and space heating drives gas and coal generation. What the heatmap makes clear is that these two extremes are not separated by chance; they are structurally determined by the interaction of Germany’s solar build-out with the seasonal daylight cycle and heating demand patterns. A practitioner optimizing scheduling on local infrastructure is therefore not searching for a random low-intensity moment; they are navigating a predictable structure where the best and worst times are largely known in advance. The summer-midday to winter-evening gradient spans the full two-to-one ratio identified earlier, and represents the maximum emissions reduction achievable through scheduling on the German grid without any hardware or algorithmic changes.

Taken together, these findings introduce a dimension of ecological efficiency that the main benchmark cannot capture: timing. The static CodeCarbon factor is not grossly wrong as an annual average, but it flattens a distribution that is wide enough to change the practical interpretation of the results. The emissions reported in this study should therefore be understood as contingent on when the runs were executed, not as intrinsic




Figure 6.10: Mean carbon intensity (gCO_2/kWh) by hour of day and calendar month for the German grid (2024), derived from 8,760 hourly observations.

properties of the algorithms. For small datasets such as Wine and Credit, this contingency is inconsequential in absolute terms, as the runs are too cheap for timing to shift the total by a meaningful amount. For large-scale workloads such as the MLP on HIGGS, grid-aware scheduling would have reduced the training footprint by the equivalent of the entire Credit dataset’s combined model training budget. At that scale, scheduling is no longer a marginal consideration but a first-order efficiency lever, comparable in impact to moderate algorithmic improvements, and achievable without touching a single line of model code.

6.8 Lifecycle Break-Even Analysis

Figure 6.11 shows for each model-dataset combination the number of predictions at which cumulative inference CO₂ surpasses the one-time training footprint.




Figure 6.11: Break-even number of predictions at which cumulative inference CO₂ emissions surpass the one-time training footprint (log scale).

Zahlen nennen: Break-even Punkte je Modell und Dataset. LR vs. MLP Kontrast besonders relevant.

Bibliography

- [1] R. Schwartz, J. Dodge, N. A. Smith, and O. Etzioni, “Green AI,” *Commun. ACM*, vol. 63, no. 12, pp. 54–63, Nov. 2020, issn: 0001-0782. doi: 10.1145/3381831 Accessed: Mar. 19, 2026.
- [2] R. Verdecchia, L. Cruz, J. Sallou, M. Lin, J. Wickenden, and E. Hotellier, “Data-Centric Green AI An Exploratory Empirical Study,” in *2022 International Conference on ICT for Sustainability (ICT4S)*, Plovdiv, Bulgaria: IEEE, Jun. 2022, pp. 35–45, isbn: 978-1-6654-8286-8. doi: 10.1109/ICT4S55073.2022.00015 Accessed: Mar. 26, 2026.
- [3] E. Strubell, A. Ganesh, and A. McCallum, *Energy and Policy Considerations for Deep Learning in NLP*, Jun. 2019. doi: 10.48550/arXiv.1906.02243 arXiv: 1906.02243 [cs]. Accessed: Mar. 19, 2026.
- [4] *Green AI Position Paper*, <https://greensoftware.foundation/policy/research/green-ai-position-paper/>. Accessed: Mar. 30, 2026.
- [5] T. Eilam, P. Bello-Maldonado, B. Bhattacharjee, C. Costa, E. K. Lee, and A. Tantawi, “Towards a Methodology and Framework for AI Sustainability Metrics,” in *Proceedings of the 2nd Workshop on Sustainable Computer Systems*, ser. HotCarbon ’23, New York, NY, USA: Association for Computing Machinery, Aug. 2023, pp. 1–7, isbn: 979-8-4007-0242-6. doi: 10.1145/3604930.3605715 Accessed: Apr. 10, 2026.
- [6] P. Henderson, J. Hu, J. Romoff, E. Brunskill, D. Jurafsky, and J. Pineau, “Towards the Systematic Reporting of the Energy and Carbon Footprints of Machine Learning,” *Journal of Machine Learning Research*, vol. 21, no. 248, pp. 1–43, 2020, issn: 1533-7928. Accessed: May 6, 2026.
- [7] A. Lacoste, A. Luccioni, V. Schmidt, and T. Dandres, *Quantifying the Carbon Emissions of Machine Learning*, Nov. 2019. doi: 10.48550/arXiv.1910.09700 arXiv: 1910.09700 [cs]. Accessed: Mar. 19, 2026.
- [8] J. Dodge et al., “Measuring the Carbon Intensity of AI in Cloud Instances,” in *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, ser. FAccT ’22, New York, NY, USA: Association for Computing Machinery, Jun. 2022, pp. 1877–1894, isbn: 978-1-4503-9352-2. doi: 10.1145/3531146.3533234 Accessed: Mar. 11, 2026.
- [9] C. Rodriguez, L. Degioanni, L. Kamení, R. Vidal, and G. Neglia, *Evaluating the Energy Consumption of Machine Learning: Systematic Literature Review and Experiments*, Aug. 2024. doi: 10.48550/arXiv.2408.15128 arXiv: 2408.15128 [cs]. Accessed: Mar. 11, 2026.

- [10] N. Bannour, S. Ghannay, A. Névél, and A.-L. Ligozat, “Evaluating the carbon footprint of NLP methods: A survey and analysis of existing tools,” in *Proceedings of the Second Workshop on Simple and Efficient Natural Language Processing*, Virtual: Association for Computational Linguistics, 2021, pp. 11–21. DOI: 10.18653/v1/2021.sustainlp-1.2 Accessed: Mar. 26, 2026.
- [11] L. Lannelongue, J. Grealey, and M. Inouye, “Green Algorithms: Quantifying the Carbon Footprint of Computation,” *Advanced Science*, vol. 8, no. 12, p. 2100707, 2021, ISSN: 2198-3844. DOI: 10.1002/advs.202100707 Accessed: May 4, 2026.
- [12] L. F. W. Anthony, B. Kanding, and R. Selvan, *Carbontracker: Tracking and Predicting the Carbon Footprint of Training Deep Learning Models*, Jul. 2020. DOI: 10.48550/arXiv.2007.03051 arXiv: 2007.03051 [cs]. Accessed: May 4, 2026.
- [13] B. Courty et al., *Mlco2/codecarbon: V2.4.1*, Zenodo, May 2024. DOI: 10.5281/zenodo.11171501 Accessed: Apr. 15, 2026.
- [14] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning*. Springer, 2017, ISBN: 978-0-387-84858-7. Accessed: May 4, 2026.
- [15] Y. Gorishniy, I. Rubachev, V. Khrulkov, and A. Babenko, “Revisiting Deep Learning Models for Tabular Data,” in *Advances in Neural Information Processing Systems*, vol. 34, Curran Associates, Inc., 2021, pp. 18 932–18 943. Accessed: Mar. 26, 2026.
- [16] N. Srivastava, G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov, “Dropout: A Simple Way to Prevent Neural Networks from Overfitting,” *Journal of Machine Learning Research*, vol. 15, no. 56, pp. 1929–1958, 2014, ISSN: 1533-7928. Accessed: Apr. 18, 2026.
- [17] J. H. Friedman, “Greedy function approximation: A gradient boosting machine,” *The Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, Oct. 2001, ISSN: 0090-5364, 2168-8966. DOI: 10.1214/aos/1013203451 Accessed: Apr. 19, 2026.
- [18] D. R. Cox, “The Regression Analysis of Binary Sequences,” *Journal of the Royal Statistical Society. Series B (Methodological)*, vol. 20, no. 2, pp. 215–242, 1958, ISSN: 0035-9246. JSTOR: 2983890. Accessed: Apr. 8, 2026.
- [19] F. Pedregosa et al., “Scikit-learn: Machine Learning in Python,” *Journal of Machine Learning Research*, vol. 12, no. 85, pp. 2825–2830, 2011, ISSN: 1533-7928. Accessed: Mar. 19, 2026.
- [20] J. R. Quinlan, “Induction of decision trees,” *Machine Learning*, vol. 1, no. 1, pp. 81–106, Mar. 1986, ISSN: 1573-0565. DOI: 10.1007/BF00116251 Accessed: Apr. 29, 2026.
- [21] L. Grinsztajn, E. Oyallon, and G. Varoquaux, “Why do tree-based models still outperform deep learning on typical tabular data?” *Advances in Neural Information Processing Systems*, vol. 35, pp. 507–520, Dec. 2022. Accessed: Apr. 29, 2026.
- [22] L. Breiman, “Random Forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, Oct. 2001, ISSN: 1573-0565. DOI: 10.1023/A:1010933404324 Accessed: Mar. 19, 2026.
- [23] L. Breiman, “Bagging predictors,” *Machine Learning*, vol. 24, no. 2, pp. 123–140, Aug. 1996, ISSN: 1573-0565. DOI: 10.1007/BF00058655 Accessed: Apr. 29, 2026.

- [24] E. Scornet, G. Biau, and J.-P. Vert, "Consistency of Random Forests," *The Annals of Statistics*, vol. 43, Apr. 2014. DOI: 10.1214/15-AOS1321
- [25] G. Biau and E. Scornet, "A random forest guided tour," *TEST*, vol. 25, no. 2, pp. 197–227, Jun. 2016, ISSN: 1863-8260. DOI: 10.1007/s11749-016-0481-7 Accessed: Apr. 29, 2026.
- [26] T. Chen and C. Guestrin, "XGBoost: A Scalable Tree Boosting System," in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, ser. KDD '16, New York, NY, USA: Association for Computing Machinery, Aug. 2016, pp. 785–794, ISBN: 978-1-4503-4232-2. DOI: 10.1145/2939672.2939785 Accessed: Mar. 19, 2026.
- [27] R. Mitchell and E. Frank, "Accelerating the XGBoost algorithm using GPU computing," *PeerJ Computer Science*, vol. 3, e127, Jul. 2017, ISSN: 2376-5992. DOI: 10.7717/peerj-cs.127 Accessed: Apr. 29, 2026.
- [28] K. Hornik, M. Stinchcombe, and H. White, "Multilayer feedforward networks are universal approximators," *Neural Networks*, vol. 2, no. 5, pp. 359–366, Jan. 1989, ISSN: 0893-6080. DOI: 10.1016/0893-6080(89)90020-8 Accessed: Apr. 29, 2026.
- [29] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representations by back-propagating errors," *Nature*, vol. 323, no. 6088, pp. 533–536, Oct. 1986, ISSN: 1476-4687. DOI: 10.1038/323533a0 Accessed: Apr. 19, 2026.
- [30] D. P. Kingma and L. J. Ba, "Adam: A Method for Stochastic Optimization," 2015. Accessed: Apr. 19, 2026.
- [31] Y. Bengio, P. Simard, and P. Frasconi, "Learning long-term dependencies with gradient descent is difficult," *IEEE Transactions on Neural Networks*, vol. 5, no. 2, pp. 157–166, Mar. 1994, ISSN: 1941-0093. DOI: 10.1109/72.279181 Accessed: Apr. 29, 2026.
- [32] G. Hinton, V. Nair, and Y. Rachmad, "Rectified Linear Units Improve Restricted Boltzmann Machines Vinod Nair," vol. 27, pp. 807–814, Jun. 2010.
- [33] S. Ioffe and C. Szegedy, "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift," in *Proceedings of the 32nd International Conference on Machine Learning*, PMLR, Jun. 2015, pp. 448–456. Accessed: Apr. 18, 2026.
- [34] L. Prechelt, "Early Stopping — But When?" In *Neural Networks: Tricks of the Trade: Second Edition*, G. Montavon, G. B. Orr, and K.-R. Müller, Eds., Berlin, Heidelberg: Springer, 2012, pp. 53–67, ISBN: 978-3-642-35289-8. DOI: 10.1007/978-3-642-35289-8_5 Accessed: Apr. 29, 2026.
- [35] J. Nickolls and W. J. Dally, "The GPU Computing Era," *IEEE Micro*, vol. 30, no. 2, pp. 56–69, Mar. 2010, ISSN: 1937-4143. DOI: 10.1109/MM.2010.41 Accessed: May 4, 2026.
- [36] P. Koshute, J. Zook, and I. McCulloh, *Recommending Training Set Sizes for Classification*, <https://arxiv.org/abs/2102.09382v1>, Feb. 2021. Accessed: Mar. 25, 2026.

- [37] A. Althnian et al., “Impact of Dataset Size on Classification Performance: An Empirical Evaluation in the Medical Domain,” *Applied Sciences*, vol. 11, no. 2, p. 796, Jan. 2021, issn: 2076-3417. doi: 10.3390/app11020796 Accessed: Apr. 16, 2026.
- [38] M. Sokolova and G. Lapalme, “A systematic analysis of performance measures for classification tasks,” *Information Processing & Management*, vol. 45, no. 4, pp. 427–437, Jul. 2009, issn: 0306-4573. doi: 10.1016/j.ipm.2009.03.002 Accessed: May 4, 2026.
- [39] R. Desislavov, F. Martínez-Plumed, and J. Hernández-Orallo, “Trends in AI inference energy consumption: Beyond the performance-vs-parameter laws of deep learning,” *Sustainable Computing: Informatics and Systems*, vol. 38, p. 100857, Apr. 2023, issn: 22105379. doi: 10.1016/j.suscom.2023.100857 Accessed: Apr. 10, 2026.
- [40] L. Yang and A. Shami, “On hyperparameter optimization of machine learning algorithms: Theory and practice,” *Neurocomputing*, vol. 415, pp. 295–316, Nov. 2020, issn: 0925-2312. doi: 10.1016/j.neucom.2020.07.061 Accessed: May 4, 2026.
- [41] J. Bergstra and Y. Bengio, “Random Search for Hyper-Parameter Optimization,” *Journal of Machine Learning Research*, vol. 13, no. 10, pp. 281–305, 2012, issn: 1533-7928. Accessed: Apr. 19, 2026.
- [42] J. Bergstra, R. Bardenet, Y. Bengio, and B. Kégl, “Algorithms for Hyper-Parameter Optimization,” in *Advances in Neural Information Processing Systems*, vol. 24, Curran Associates, Inc., 2011. Accessed: Apr. 19, 2026.
- [43] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, “Optuna: A Next-generation Hyperparameter Optimization Framework,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, ser. KDD ’19, New York, NY, USA: Association for Computing Machinery, Jul. 2019, pp. 2623–2631, isbn: 978-1-4503-6201-6. doi: 10.1145/3292500.3330701 Accessed: Apr. 1, 2026.
- [44] R. Fischer, *Ground-Truthing AI Energy Consumption: Validating CodeCarbon Against External Measurements*, Sep. 2025. doi: 10.48550/arXiv.2509.22092 arXiv: 2509.22092 [cs]. Accessed: May 6, 2026.
- [45] D. Patterson et al., *Carbon Emissions and Large Neural Network Training*, Apr. 2021. doi: 10.48550/arXiv.2104.10350 arXiv: 2104.10350 [cs]. Accessed: May 9, 2026.
- [46] C.-J. Wu et al., “Sustainable AI: Environmental Implications, Challenges and Opportunities,” *Proceedings of Machine Learning and Systems*, vol. 4, pp. 795–813, Apr. 2022. Accessed: May 9, 2026.
- [47] S. Luccioni, Y. Jernite, and E. Strubell, “Power Hungry Processing: Watts Driving the Cost of AI Deployment?” In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, ser. FAccT ’24, New York, NY, USA: Association for Computing Machinery, Jun. 2024, pp. 85–99, isbn: 979-8-4007-0450-5. doi: 10.1145/3630106.3658542 Accessed: May 9, 2026.
- [48] D. Hernandez and T. B. Brown, *Measuring the Algorithmic Efficiency of Neural Networks*, May 2020. doi: 10.48550/arXiv.2005.04305 arXiv: 2005.04305 [cs]. Accessed: May 9, 2026.

-
- [49] C. Coleman et al., “Analysis of DAWN Bench, a Time-to-Accuracy Machine Learning Performance Benchmark,” *SIGOPS Oper. Syst. Rev.*, vol. 53, no. 1, pp. 14–25, Jul. 2019, issn: 0163-5980. doi: 10.1145/3352020.3352024 Accessed: May 9, 2026.
- [50] P. Mattson et al., “MLPerf Training Benchmark,” *Proceedings of Machine Learning and Systems*, vol. 2, pp. 336–349, Mar. 2020. Accessed: May 9, 2026.
- [51] M. F. Stefan Aeberhard, *Wine*, 1992. doi: 10.24432/C5PC7J Accessed: Mar. 19, 2026.
- [52] I-Cheng Yeh, *Default of Credit Card Clients*, 2009. doi: 10.24432/C55S3H Accessed: Mar. 19, 2026.
- [53] P. Baldi, P. Sadowski, and D. Whiteson, “Searching for Exotic Particles in High-Energy Physics with Deep Learning,” *Nature Communications*, vol. 5, no. 1, p. 4308, Jul. 2014, issn: 2041-1723. doi: 10.1038/ncomms5308 arXiv: 1402.4735 [hep-ph]. Accessed: Mar. 19, 2026.
- [54] D. Whiteson, *HIGGS*, 2014. doi: 10.24432/C5V312 Accessed: Mar. 19, 2026.
- [55] A. Bagnall and G. C. Cawley, *On the Use of Default Parameter Settings in the Empirical Evaluation of Classification Algorithms*, Mar. 2017. doi: 10.48550/arXiv.1703.06777 arXiv: 1703.06777 [cs]. Accessed: Apr. 10, 2026.
- [56] A. Paszke et al., *PyTorch: An Imperative Style, High-Performance Deep Learning Library*, Dec. 2019. doi: 10.48550/arXiv.1912.01703 arXiv: 1912.01703 [cs]. Accessed: Mar. 19, 2026.
- [57] B. Ghosh, *Skorch: The Bridge Between PyTorch and Scikit-Learn*, Sep. 2024. Accessed: Mar. 26, 2026.
- [58] *Claude Code*, Anthropic, 2026.
- [59] N. Feix, *Snip3rnick/PyHardwareMonitor*, Mar. 2026. Accessed: May 8, 2026.
- [60] *Electricity Maps API: Carbon Intensity History*. Accessed: Apr. 20, 2026.

List of Figures

2.1	Random Forest vs Gradient Boosting: Parallel (breadth) vs Sequential (depth) ensemble strategies.	12
5.1	Visualization of the data loading pipeline and preprocessing steps within the <code>load_data()</code> method.	42
5.2	Training and emissions tracking procedure per model-dataset combination. Dashed gray boxes indicate CodeCarbon <code>EmissionsTracker</code> calls, and dashed blue boxes indicate <code>CPUPowerMonitor</code> calls (background thread, 0.25 s sampling). Both run in parallel during cross-validation. After training, <code>compute_corrected_co2</code> replaces CodeCarbon’s TDP-based CPU estimate with the hardware sensor value.	48
6.1	CodeCarbon underestimation factor (hardware-corrected / CodeCarbon) per model and dataset. All values exceed 1, confirming systematic undercounting. Higher factors for Wine and Credit reflect CodeCarbon’s TDP-based CPU estimation being least accurate on short training runs; RF has no HIGGS measurement.	57
6.2	Model comparison dashboard: CO ₂ emissions (corrected), training time, inference time, accuracy, F1-score, and carbon-penalised F1 (EWF1) across all models and datasets.	59
6.3	XGBoost CPU vs. GPU across datasets. CO ₂ and training time on log scale. .	60
6.4	XGBoost break-even point comparison between different subset sizes of the HIGGS dataset.	61
6.5	All models evaluated on different scales of the HIGGS subsets.	61
6.6	Carbon intensity (gCO ₂ /kWh) for eight timing scenarios	64
6.7	Counterfactual CO ₂ relative to the static CodeCarbon baseline (= 100 %) . . .	65
6.8	Hourly carbon intensity on 24 July 2024 in Germany	65
6.9	Hourly carbon intensity for representative ten-day windows per season, Germany 2025	66
6.10	Mean carbon intensity (gCO ₂ /kWh) by hour of day and calendar month for the German grid (2024), derived from 8,760 hourly observations.	67
6.11	Break-even number of predictions at which cumulative inference CO ₂ emissions surpass the one-time training footprint (log scale).	68

List of Tables

2.1	Binary confusion matrix with the four fundamental prediction outcomes. . .	18
2.2	Comparison of HPO strategies (n : trials, d : hyperparameters).	23
4.1	Overview of datasets used in this analysis	36
4.2	Overview of classification models used in this study	37
5.1	Hardware specification of the experimental system	39
5.2	Optimized hyperparameters for the Random Forest classifier obtained via Optuna optimization.	43
5.3	Optimized hyperparameters for the XGBoost classifier across all datasets. . .	44
5.4	Resolved layer-wise architecture of the tuned MLP for each dataset. Each hidden block consists of a Linear layer, Batch Normalization, ReLU, and Dropout. Parameter counts include weights, biases, and Batch Normalization scale and shift parameters.	45
5.5	Hyperparameter search spaces defined in the Optuna tuning scripts.	46
5.6	Column schema of <code>results.csv</code> including component-level energy tracking.	51
5.7	Column schema of <code>inference_time.csv</code>	51
5.8	MLP architectures evaluated in the variation experiment, with trainable parameter counts for HIGGS. Counts include linear layer weights and biases.	53
5.9	Carbon intensity scenarios used in the counterfactual analysis. Each scenario is computed from one year of hourly ElectricityMaps data for the German bidding zone.	55
6.1	Hyperparameter tuning overhead per model and dataset. Each run comprised 40 Optuna trials optimising weighted F_1 . CO_2 was measured with the hardware-corrected tracker; duration covers the full <code>study.optimize</code> call. Random Forest was not tuned on HIGGS; literature defaults were used instead.	58