



Karlsruhe University of Applied Sciences

Department of Computer Science and Business Information Systems

Business Information Systems

BACHELOR'S THESIS

Efficiency vs. Performance in Green AI: An Empirical Investigation of Energy Consumption, Carbon Footprint, and Predictive Accuracy of Classification Algorithms at Varying Dataset Sizes.

Author	Eron Qorri
Matriculation number	85383
Workplace	Hochschule Karlsruhe
First Advisor	Prof. Dr. rer. nat. Christine Preisach
Second Advisor	Prof. Dr. Reimar Hofmann
Closing Date	31 July, 2026

31 July, 2026

Chairman of the Examination Board

Contents

1	Introduction	1
1.1	Problem statement and societal relevance	1
1.2	Motivation: when does a simpler model pay off?	2
1.3	Research questions and objectives	2
1.4	Structure of the Thesis	2
	Bibliography	5
	List of Figures	6
	List of Tables	7

Chapter 1

Introduction

1.1 Problem statement and societal relevance

The current trend in machine learning is moving towards increasingly large models and datasets [1]. Although deep learning approaches often achieve the highest predictive accuracy, their energy demand and associated CO_2 emissions increase disproportionately [2], [3]. In an era marked by strict sustainability guidelines and rising energy costs, a critical question arises: At what point is a computationally "simple" model ecologically more viable than a highly complex one?

The primary objective of this thesis is to determine whether and when it is worthwhile to deploy more complex models, considering not only their predictive accuracy but also their associated energy consumption. To achieve this, a Logistic Regression, a Random Forest, an XGBoost model and a Multilayer Perceptron are compared under equal conditions across three datasets of varying sizes. This comparative analysis aims to provide a clearer understanding of the thresholds at which simpler models represent the better choice. Additionally, the study investigates how modifying an Multilayer Perceptron architecture, specifically varying the number of neurons and hidden layers, directly impacts energy consumption and CO_2 emissions.

To address the problem statement and achieve the outlined objectives, this thesis investigates the overarching question: *How does the relationship between energy consumption and predictive accuracy change for Random Forest, XGBoost, and Multilayer Perceptron when model complexity and dataset size are systematically varied?*

Specifically, the following sub-questions are examined:

- How do energy consumption (in kWh) and CO_2 emissions change when scaling a deep learning architecture (variation of neuron count)?
- Under which conditions (dataset size, model complexity) does the resource consumption of complex models exceed the achievable accuracy gain compared to simpler algorithms?
- How does the energy consumption differ between CPU-based training (Random Forest) and GPU-accelerated training (XGBoost, MLP) for the same classification task?

1.2 Motivation: when does a simpler model pay off?

1.3 Research questions and objectives

1.4 Structure of the Thesis

The remainder of this thesis is structured as follows: Chapter 2 provides the theoretical background, distinguishing between Green AI and Red AI, and introducing the relevant sustainability and classification metrics. Chapter 3 details the methodology, including dataset selection and the experimental hardware setup. Chapter 4 outlines the implementation process, focusing on data preprocessing, model training, and the logging of resource consumption. The empirical results are analyzed in Chapter 5, evaluating both performance and resource utilization across scaling dataset sizes. Chapter 6 discusses the findings, highlighting the economic and environmental implications. Finally, Chapter 7 concludes the thesis and provides actionable recommendations for sustainable machine learning practices.

Table 1: Data Set Summary

Classification Data Set	Code	# Data Points (N_D)	# Features (N_F)
Annealing [20]	ANN	798	38
Breast Cancer [21]	BCW	683	9
Bank Marketing [22]	BMK	41188	20
Cardiotocography [23]	CAR	2126	21
Car Evaluation [24]	CEV	1728	6
Congressional Voting Records [25]	CVR	435	16
Credit Card Defaults [26]	DCC	30000	23
Diabetic Retinopathy [27]	DRB	1151	19
Handwritten Digit Recognition [28]	DRE	5619	64
Indian Liver Patient [29]	ILP	583	10
Ionosphere [30]	ION	351	34
Letter Recognition [31]	LRE	20000	16
Gamma Telescope Particles [32]	MAG	19020	10
Mice Protein Expression [33]	MPE	972	80
Mushroom Edibility [34]	MSH	8124	22
Occupancy Detection [35]	OCD	20560	5
Phishing Websites [36]	PHW	11044	30
QSAR Biodegradability [37]	QBD	1055	41
Skin Segmentation [38]	SSG	245057	3
Wall-Following Robots [39]	WFR	5456	24

I could do something of that sort for my own data sets

We also gathered several other characteristics of each data set, including number of categorical features (N_{CAT}), number of continuous features (N_{CONT}), proportion of features that are categorical (R_{CAT}), number of target classes (N_C), minimum single-class proportion (C_{MIN}), and class imbalance (I_C). The latter is a measure of the difference in actual and expected proportion of the least common target class (L) and its expected value (LE); i.e., $I = (LE - L)/LE$. These values are shown in Table 2.

Table 2: Additional Data Set Characteristics

Set	N_{CAT}	N_{CONT}	R_{CAT}	N_C	C_{MIN}	I_C
ANN	32	6	0.842	5	0.001	0.956
BCW	0	9	0.000	2	0.350	0.300
BMK	10	10	0.500	2	0.113	0.775
CAR	0	21	0.000	3	0.083	0.752
CVE	6	0	1.000	4	0.038	0.850
CVR	16	0	1.000	2	0.386	0.228
DCC	3	20	0.130	2	0.221	0.558
DRB	0	19	0.000	2	0.469	0.062
DRE	0	64	0.000	10	0.098	0.016
ILP	1	9	0.100	2	0.286	0.427
ION	0	34	0.000	2	0.359	0.282
LRE	0	16	0.000	26	0.037	0.046
MAG	0	10	0.000	2	0.322	0.357
MPE	3	77	0.038	8	0.092	0.259
MSH	22	0	1.000	2	0.482	0.036
OCD	0	5	0.000	2	0.231	0.538
PHW	30	0	1.000	2	0.443	0.114
QBD	0	41	0.000	2	0.337	0.325
SSG	0	3	0.000	2	0.208	0.585
WFR	0	24	0.000	4	0.060	0.760

Bibliography

- [1] R. Schwartz, J. Dodge, N. A. Smith, and O. Etzioni, *Green AI*, Aug. 13, 2019. DOI: 10.48550/arXiv.1907.10597 arXiv: 1907.10597[cs]. Accessed: Mar. 19, 2026. [Online]. Available: <http://arxiv.org/abs/1907.10597>
- [2] E. Strubell, A. Ganesh, and A. McCallum, *Energy and policy considerations for deep learning in NLP*, Jun. 5, 2019. DOI: 10.48550/arXiv.1906.02243 arXiv: 1906.02243[cs]. Accessed: Mar. 19, 2026. [Online]. Available: <http://arxiv.org/abs/1906.02243>
- [3] A. Lacoste, A. Luccioni, V. Schmidt, and T. Dandres, *Quantifying the carbon emissions of machine learning*, Nov. 4, 2019. DOI: 10.48550/arXiv.1910.09700 arXiv: 1910.09700[cs]. Accessed: Mar. 19, 2026. [Online]. Available: <http://arxiv.org/abs/1910.09700>

List of Figures

List of Tables